

Introduction

The Nasdaq-100 is a stock market index that includes a selection of 102 equity securities issued by 101 of the largest nonfinancial companies listed on the Nasdaq stock exchange. This index represents a diverse range of sectors, including manufacturing, technology, retail, telecommunication, biotechnology, health care, transportation, media, and service providers.

Key Points about the Nasdaq-100 Index

- **Composition:** The Nasdaq-100 is comprised of 102 stocks from various sectors, with a significant focus on technology-related companies. It does not include financial companies like banks and insurance firms.
- **Market Capitalization:** The index is weighted by market capitalization, meaning that larger companies have a more significant impact on the index's value.
- **Tech Heavy:** It is often considered a tech-heavy index due to the inclusion of prominent technology companies like Apple, Amazon, Microsoft, and others.
- **Diversity:** While it has a technology emphasis, the Nasdaq-100 is diversified across different sectors, which can reduce risks associated with investing in a single industry.
- **Global Reach:** Many of the companies in the Nasdaq-100 have a global presence and are leaders in their respective industries.
- **Investment and Trading:** The Nasdaq-100 is a popular choice for investors and traders looking to gain exposure to a broad range of innovative and growth-oriented companies.
- **Volatility:** Given its composition of technology and high-growth stocks, the Nasdaq-100 can experience higher volatility compared to other indices like the S&P 500.
- **Benchmark:** It is often used as a benchmark for the performance of technology and growth stocks.

Investors and financial professionals use the Nasdaq-100 as a reference point to gauge the performance of the technology sector and the broader market. It is also a basis for various financial products, including ETFs (Exchange-Traded Funds) that track its performance.

Section 1: Applied Data Science With Python

Question1

Check the stock symbols of the companies in `Nasdaq 100 Market cap.xlsx`. Only the relevant files in the `NASDAQ_DATA` folder should be read.

```

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.holtwinters import ExponentialSmoothing
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
import pmdarima as pm
from pmdarima.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error

```

Question 2

Append all files (imported in the previous step) that contain no more than 10 years of data. For this, you may use your discretion.

Question 3

Read `Nasdaq 100 market cap.xlsx` and `nasdaq100_metrics_ratios.xlsx`.

```

metrics_ratio=pd.read_excel(r"D:\Share Market Analysis Dataset\
nasdaq100_metrics_ratios.xlsx")
market_cap=pd.read_excel(r"D:\Share Market Analysis Dataset\Nasdaq 100
Market cap.xlsx")

metrics_ratio.shape

(102, 283)

market_cap.shape

(102, 6)

#### Perform the join
result = pd.merge(metrics_ratio, market_cap, left_on='symbol',
right_on='Symbol')

#### Drop the duplicate 'Symbol' column
result = result.drop('Symbol', axis=1)

result.shape

(102, 288)

strings = ['AAPL', 'ABNB', 'ADBE', 'ADI', 'ADP', 'ADSK', 'AEP',
'ALGN',

```

```

'AMAT', 'AMD', 'AMGN', 'AMZN', 'ANSS', 'ASML', 'ATVI', 'AVGO',
'AZN', 'BIDU', 'BIIB', 'BKNG', 'CDNS', 'CEG', 'CHTR', 'CMCSA',
'COST', 'CPRT', 'CRWD', 'CSCO', 'CSX', 'CTAS', 'CTSH', 'DDOG',
'DLTR', 'DOCU', 'DXCM', 'EA', 'EBAY', 'EXC', 'FAST', 'FISV',
'FTNT', 'GILD', 'GOOG', 'GOOGL', 'HON', 'IDXX', 'ILMN', 'INTC',
'INTU', 'ISRG', 'JD', 'KDP', 'KHC', 'KLAC', 'LCID', 'LRCX',
'LULU',
'MAR', 'MCHP', 'MDLZ', 'MELI', 'META', 'MNST', 'MRNA', 'MRVL',
'MSFT', 'MTCH', 'MU', 'NFLX', 'NTES', 'NVDA', 'NXPI', 'ODFL',
'OKTA', 'ORLY', 'PANW', 'PAYX', 'PCAR', 'PDD', 'PEP', 'PYPL',
'QCOM', 'REGN', 'ROST', 'SBUX', 'SGEN', 'SIRI', 'SNPS', 'SPLK',
'SWKS', 'TEAM', 'TMUS', 'TSLA', 'TXN', 'VRSK', 'VRSN', 'VRTX',
'WBA', 'WDAY', 'XEL', 'ZM', 'ZS']

import glob
file_names=glob.glob(r"D:\Share Market Analysis Dataset\NASDAQ_DATA\
*.csv")

matching_files = [file for file in file_names for string in strings if
string in file]

len(matching_files)

117

import os

```

Question 4

Collate the two files imported in the previous step to include the fields Market cap and Last sale in addition to the various metrics and ratios already present in nasdaq100_metrics_ratios.xlsx.

```

#### Initialize an empty DataFrame
all_data = pd.DataFrame()

#### Iterate over each file path
for filepath in matching_files:
    #### Read the data from the file
    data = pd.read_csv(filepath)

    #### Get the filename without extension
    filename_without_extension =
os.path.splitext(os.path.basename(filepath))[0]

    #### Add filename as a new column
    data['filename'] = filename_without_extension

    #### Append the data to the all_data DataFrame
    all_data = pd.concat([all_data, data], ignore_index=True)

```

```

all_data.columns
Index(['Date', 'High', 'Low', 'Open', 'Close', 'Volume', 'Adj Close',
      'filename', 'Adjusted Close'],
      dtype='object')

result.columns
Index(['symbol', 'company', 'sector', 'subsector',
      'asset_turnover_2017',
      'asset_turnover_2018', 'asset_turnover_2019',
      'asset_turnover_2020',
      'asset_turnover_2021', 'asset_turnover_2022',
      ...
      'yoy_revenue_growth_2019', 'yoy_revenue_growth_2020',
      'yoy_revenue_growth_2021', 'yoy_revenue_growth_2022',
      'yoy_revenue_growth_latest', 'Name', 'Market Cap', 'Last Sale',
      'Net Change', 'Percentage Change'],
      dtype='object', length=288)

merged_data = pd.merge(result, all_data, left_on='symbol',
                        right_on='filename', how='inner')

merged_data.shape, result.shape
((484476, 297), (102, 288))

merged_data.drop_duplicates(inplace=True)

merged_data.shape
(474247, 297)

```

Question 5

Identify the variables whose variance is less than .005 (as these do not contribute to model building), and eliminate those variables.

```

filtered_data=merged_data
# Filter the dataframe to include only numerical columns
numerical_cols = filtered_data.select_dtypes(include='number').columns

# Calculate variance for each numerical variable
variance = filtered_data[numerical_cols].var()

# Identify variables with variance less than 0.005
low_variance_cols = variance[variance < 0.005].index

# Drop low variance variables
filtered_data = filtered_data.drop(low_variance_cols, axis=1)

```

```
low_variance_cols
Index(['capex_to_revenue_2022', 'inventory_to_revenue_2017',
      'inventory_to_revenue_2018', 'inventory_to_revenue_2022',
      'Percentage Change'],
      dtype='object')
```

We removed the columns 'capex_to_revenue_2022', 'inventory_to_revenue_2017', 'inventory_to_revenue_2018', 'inventory_to_revenue_2022', and 'Percentage Change' from the dataset because they had low variance, meaning their values didn't vary much and didn't provide much useful information for analysis.

```
#### Convert 'Date' to datetime format
merged_data['Date'] = pd.to_datetime(merged_data['Date'])

#### Get the latest date in the dataset
latest_date = merged_data['Date'].max()

#### Calculate the date 10 years ago from the latest date
date_10_years_ago = latest_date - pd.DateOffset(years=10)

#### Filter the data for the last 10 years
filtered_data = merged_data[merged_data['Date'] > date_10_years_ago]

filtered_data.shape
(236022, 297)

filtered_data.to_csv("test.csv")

filtered_data.shape
(236022, 297)
```

Question 6

Delete the variables in `nasdaq100_metrics_ratios.xlsx` where 30% or more of the values are missing.

Question 7

Perform missing value imputation for variables with less than 30% missing values by considering the company's sector. `secto`

```
#### Calculate the percentage of missing values for each variable
missing_percent = filtered_data.isnull().mean() * 100
```

```

#### Identify variables with 30% or more missing values
cols_to_drop = missing_percent[missing_percent >= 30].index

#### Drop these variables
filtered_data = filtered_data.drop(cols_to_drop, axis=1)

#### Identify variables with less than 30% missing values
cols_to_impute = missing_percent[(missing_percent > 0) &
(missing_percent < 30)].index

#### Impute missing values based on the company's sector
for col in cols_to_impute:
    filtered_data[col] = filtered_data.groupby('sector')
[col].transform(lambda x: x.fillna(x.mean()))

filtered_data.shape

(236022, 227)

```

Question 8

Analyze the effect of COVID on stock prices in detail, create visuals to support the insights, and address the following:

8.a

Which sectors and companies saw the greatest impact, and which ones saw the least? You may use growth or degrowth as a measure of impact and may perform week over week, month over month (MoM), quarter over quarter (QoQ), or year over year (YoY) analysis as appropriate.

8.b

Which sector and company experienced the fastest and slowest recoveries?

```

#### Convert the date column to datetime format
filtered_data['Date'] = pd.to_datetime(filtered_data['Date'])

#### Define the COVID impact and recovery periods
covid_start = '2020-02-01'
covid_end = '2020-06-30'
recovery_start = '2020-07-01'
recovery_end = '2021-12-31'

#### Filter data for the COVID impact period
covid_data = filtered_data[(filtered_data['Date'] >= covid_start) &
(filtered_data['Date'] <= covid_end)]

```

```

#### Calculate the growth or degrowth in stock prices during the COVID
impact period
covid_data['price_change'] = covid_data.groupby(['symbol', 'sector'])
['Close'].pct_change()

#### Aggregate the data to get the average growth or degrowth for each
sector and company
covid_agg = covid_data.groupby(['sector', 'symbol'])
['price_change'].mean().reset_index()

#### Identify the sectors and companies with the greatest and least
impact
greatest_impact_sector = covid_agg.groupby('sector')
['price_change'].mean().idxmin()
least_impact_sector = covid_agg.groupby('sector')
['price_change'].mean().idxmax()
greatest_impact_company = covid_agg['price_change'].idxmin()
least_impact_company = covid_agg['price_change'].idxmax()

#### Filter data for the recovery period
recovery_data = filtered_data[(filtered_data['Date'] >=
recovery_start) & (filtered_data['Date'] <= recovery_end)]

#### Calculate the growth or recovery in stock prices during the
recovery period
recovery_data['price_change'] = recovery_data.groupby(['symbol',
'sector'])['Close'].pct_change()

#### Aggregate the data to get the average growth or recovery for each
sector and company
recovery_agg = recovery_data.groupby(['sector', 'symbol'])
['price_change'].mean().reset_index()

#### Identify the sectors and companies with the fastest and slowest
recovery
fastest_recovery_sector = recovery_agg.groupby('sector')
['price_change'].mean().idxmax()
slowest_recovery_sector = recovery_agg.groupby('sector')
['price_change'].mean().idxmin()
fastest_recovery_company = recovery_agg['price_change'].idxmax()
slowest_recovery_company = recovery_agg['price_change'].idxmin()

#### Create visuals to support the insights
fig, ax = plt.subplots(2, 1, figsize=(10, 15))

#### Bar plot for impact on sectors
covid_agg.groupby('sector')['price_change'].mean().plot(kind='bar',
ax=ax[0])
ax[0].set_title('Impact of COVID on Stock Prices by Sector')
ax[0].set_ylabel('Average Growth or Degrowth')

```

```
#### Bar plot for recovery of sectors
recovery_agg.groupby('sector')['price_change'].mean().plot(kind='bar',
ax=ax[1])
ax[1].set_title('Recovery of Stock Prices by Sector')
ax[1].set_ylabel('Average Growth or Recovery')

plt.tight_layout()
plt.show()
```

C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\200252812.py:14:
PerformanceWarning: DataFrame is highly fragmented. This is usually
the result of calling `frame.insert` many times, which has poor
performance. Consider joining all columns at once using
pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe
= frame.copy()`

```
covid_data['price_change'] = covid_data.groupby(['symbol',
'sector'])['Close'].pct_change()
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\200252812.py:14:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

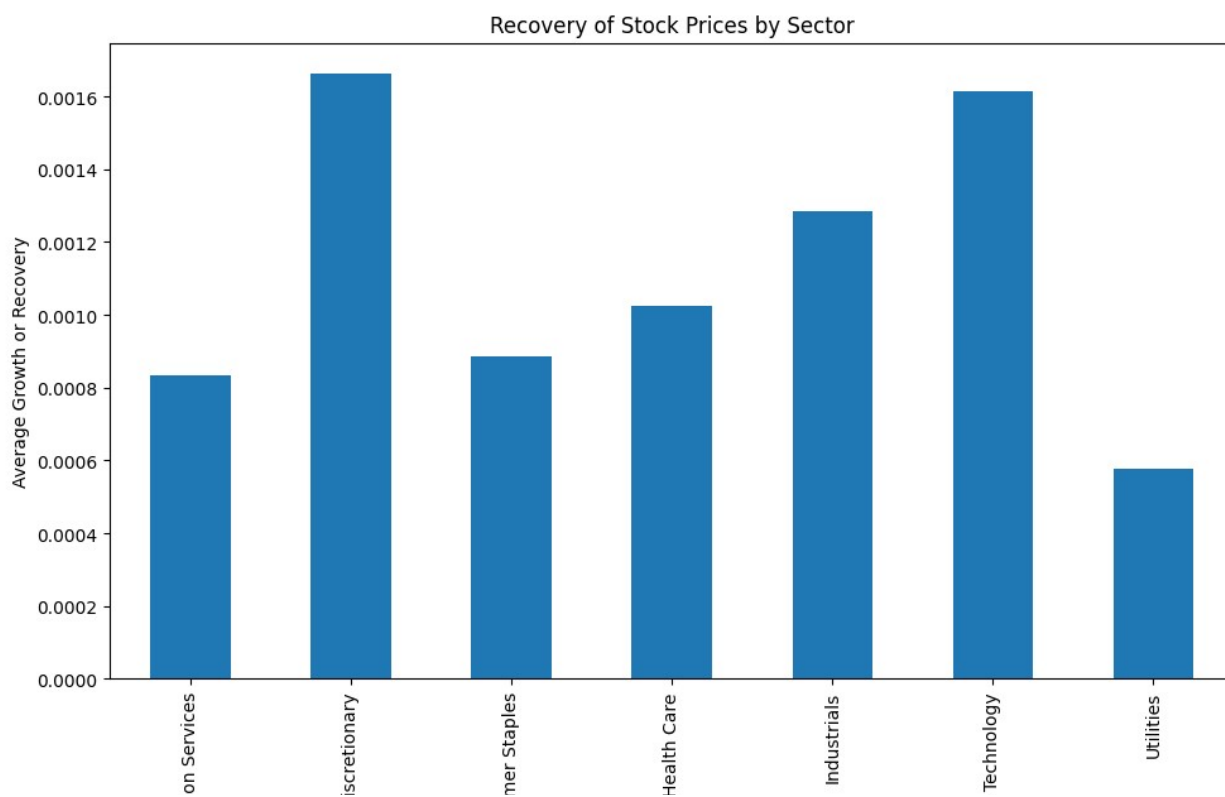
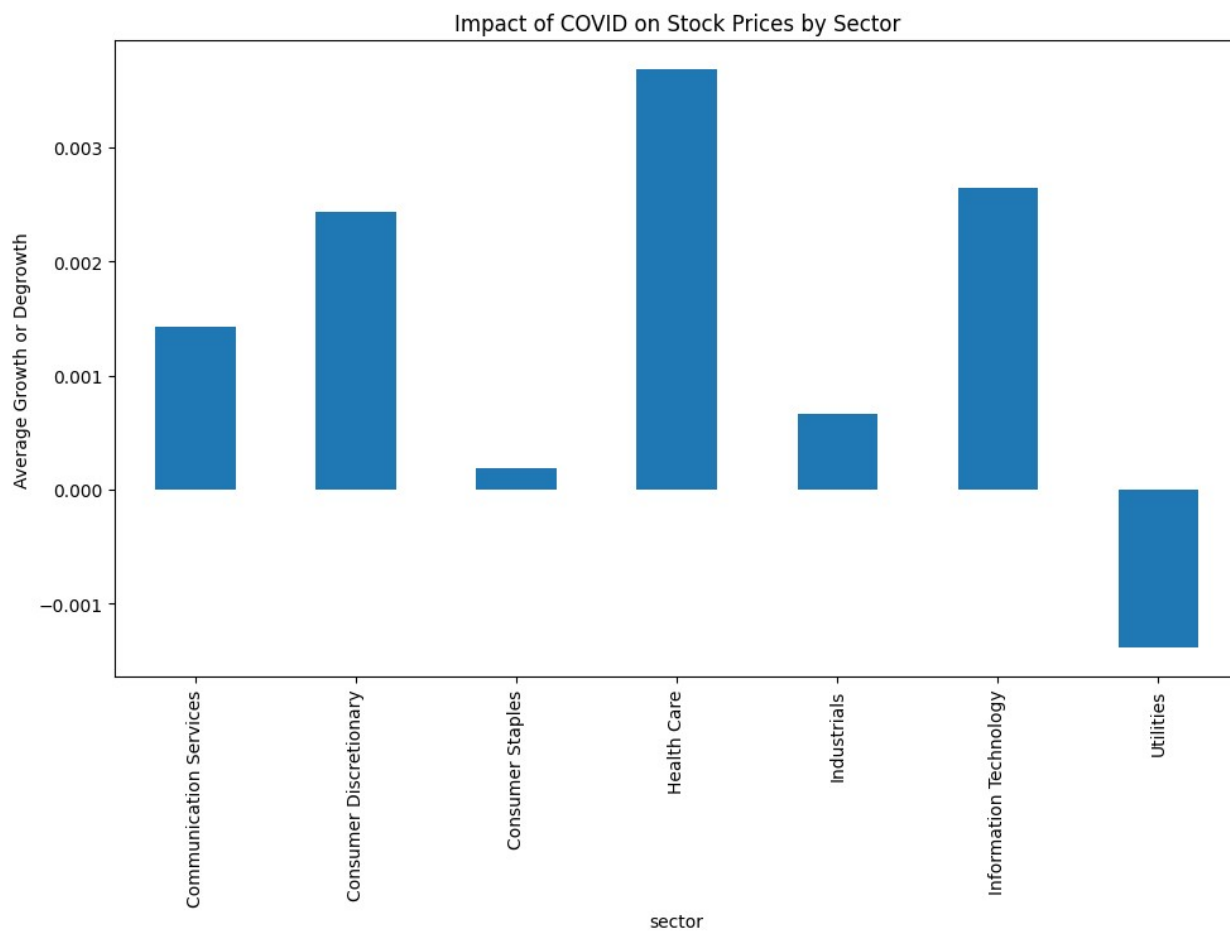
See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
covid_data['price_change'] = covid_data.groupby(['symbol',
'sector'])['Close'].pct_change()
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\200252812.py:29:
PerformanceWarning: DataFrame is highly fragmented. This is usually
the result of calling `frame.insert` many times, which has poor
performance. Consider joining all columns at once using
pd.concat(axis=1) instead. To get a de-fragmented frame, use `newframe
= frame.copy()`
```

```
recovery_data['price_change'] = recovery_data.groupby(['symbol',
'sector'])['Close'].pct_change()
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\200252812.py:29:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
recovery_data['price_change'] = recovery_data.groupby(['symbol',
'sector'])['Close'].pct_change()
```

```
filtered_data.head()
```

	symbol	company	sector	\
3228	AAPL	Apple Inc.	Information Technology	
3229	AAPL	Apple Inc.	Information Technology	
3230	AAPL	Apple Inc.	Information Technology	
3231	AAPL	Apple Inc.	Information Technology	
3232	AAPL	Apple Inc.	Information Technology	

	subsector	asset_turnover_2017
3228	Technology Hardware, Storage & Peripherals	0.66
3229	Technology Hardware, Storage & Peripherals	0.66
3230	Technology Hardware, Storage & Peripherals	0.66
3231	Technology Hardware, Storage & Peripherals	0.66
3232	Technology Hardware, Storage & Peripherals	0.66

	asset_turnover_2018	asset_turnover_2019	asset_turnover_2020	\
3228	0.72	0.74	0.83	
3229	0.72	0.74	0.83	
3230	0.72	0.74	0.83	
3231	0.72	0.74	0.83	
3232	0.72	0.74	0.83	

	asset_turnover_2021	asset_turnover_latest	...	Net Change	\
3228	1.08	0.24	...	\$2.00	
3229	1.08	0.24	...	\$2.00	
3230	1.08	0.24	...	\$2.00	
3231	1.08	0.24	...	\$2.00	
3232	1.08	0.24	...	\$2.00	

	Percentage Change	Date	High	Low	Open	\
3228	0.0134	2012-10-31	21.498571	20.989286	21.245714	
3229	0.0134	2012-11-01	21.535713	21.220358	21.365000	
3230	0.0134	2012-11-02	21.319643	20.526787	21.281786	
3231	0.0134	2012-11-05	20.991785	20.628571	20.840000	
3232	0.0134	2012-11-06	21.097857	20.717501	21.079643	

	Close	Volume	Adj Close	filename
3228	21.261429	510003200.0	18.201437	AAPL
3229	21.305000	361298000.0	18.238741	AAPL
3230	20.600000	599373600.0	17.635199	AAPL
3231	20.879286	529135600.0	17.874296	AAPL
3232	20.816071	374917200.0	17.820169	AAPL

```
[5 rows x 227 columns]
```

Section 2: Machine Learning

1 Perform PCA to reduce the number of variables in the data.

```
# Assuming you have a DataFrame named 'filtered_data' with the
specified columns

# Convert the 'Close' column to string, remove non-numeric characters,
and convert to float
filtered_data['Close'] = filtered_data['Close'].astype(str)
filtered_data['Close'] = filtered_data['Close'].str.replace('$', '',
regex=True)
filtered_data['Close'] = pd.to_numeric(filtered_data['Close'],
errors='coerce')
import pandas as pd

# Assuming you have a DataFrame named 'filtered_data'

# Loop through all columns in the DataFrame
for column in filtered_data.columns:
    # Check if the column contains a dollar symbol ($) and if it's not
    already a numeric type
    if '$' in filtered_data[column].astype(str).str.extract('(\$\d+\.\d+)'')[0].values:
        # Remove dollar symbols and convert the column to float
        filtered_data[column] = filtered_data[column].str.replace('$',
'', regex=False).astype(float)
# List of columns to process
columns_to_process = ['Last Sale', 'Net Change']

filtered_data = filtered_data.replace('[\$,]', '', regex=True)
for col in filtered_data.columns:
    if filtered_data[col].dtype == 'object':
        # If column is of object type, fill missing values with mode
        mode_val = filtered_data[col].mode()[0]
        filtered_data[col].fillna(mode_val, inplace=True)
    else:
        # If column is of numerical type, fill missing values with
        median
        median_val = filtered_data[col].median()
        filtered_data[col].fillna(median_val, inplace=True)
# Loop through columns and remove dollar symbols and commas
for column in columns_to_process:
    filtered_data[column] = filtered_data[column].str.replace('$',
'').str.replace(',', '').astype(float)

# Now, the specified columns should have dollar symbols and commas
removed and be converted to float
```

```

# Drop non-numeric columns
numerical_data = filtered_data.drop(columns=['symbol', 'company',
'sector', 'subsector', 'Name', 'filename', 'Date'])

# Handle missing values by filling them with the mean value of each
column
numerical_data.fillna(numerical_data.mean(), inplace=True)

# Standardize the data (mean=0, std=1)
scaler = StandardScaler()
data_std = scaler.fit_transform(numerical_data)

# Apply PCA with no component number restriction
pca_all = PCA()
data_pca_all = pca_all.fit_transform(data_std)

# Find the number of components required to retain 95% of the variance
cumulative_variance_ratio = pca_all.explained_variance_ratio_.cumsum()
n_components_95 = (cumulative_variance_ratio >= 0.95).argmax() + 1

# Apply PCA with the selected number of components
pca = PCA(n_components=n_components_95)
data_pca = pca.fit_transform(data_std)

# Create a DataFrame with the PCA results
data_pca_df = pd.DataFrame(data=data_pca, columns=[f'PC{i}' for i in
range(1, n_components_95 + 1)])

# Print the number of components retained
print('Number of components retained to retain 95% of the variance:',
n_components_95)

# Print the variance ratio of each component
print('Explained variance ratio:', pca.explained_variance_ratio_)

# Print the cumulative explained variance by the components
print('Cumulative explained variance:',
pca.explained_variance_ratio_.cumsum())

Number of components retained to retain 95% of the variance: 55
Explained variance ratio: [0.1330012  0.07799619 0.05904312 0.05291784
0.04514023 0.04022754
0.03567015 0.03029314 0.02806933 0.02732899 0.023224   0.02226343
0.02099559 0.02002561 0.01834272 0.01798933 0.01756261 0.01578869
0.01522894 0.01391463 0.01277926 0.01206117 0.01191486 0.01165498
0.01041486 0.00965039 0.00948009 0.00935483 0.00884144 0.00865445
0.00841754 0.00822238 0.00786466 0.00688673 0.00679613 0.00646869
0.00642051 0.0061495  0.00590385 0.00551129 0.00523013 0.00518109
0.00502464 0.00492235 0.00464825 0.00450944 0.00445227 0.00424488]

```

```

0.00394151 0.00374964 0.00371565 0.00338303 0.00332605 0.00310619
0.00304254]
Cumulative explained variance: [0.1330012 0.21099738 0.2700405
0.32295834 0.36809857 0.40832611
0.44399627 0.4742894 0.50235873 0.52968772 0.55291172 0.57517515
0.59617074 0.61619634 0.63453906 0.65252839 0.670091 0.68587969
0.70110863 0.71502326 0.72780252 0.73986369 0.75177855 0.76343353
0.77384839 0.78349878 0.79297887 0.8023337 0.81117514 0.81982959
0.82824713 0.8364695 0.84433416 0.85122089 0.85801702 0.86448571
0.87090622 0.87705572 0.88295957 0.88847086 0.89370099 0.89888208
0.90390673 0.90882907 0.91347733 0.91798676 0.92243903 0.92668392
0.93062542 0.93437506 0.93809071 0.94147374 0.94479979 0.94790598
0.95094851]

```

Summary

- Number of components retained to retain 95% of the variance: 55
- Explained variance ratio for each of the 55 components:
 - These values represent the proportion of variance explained by each component. The higher the value, the more variance is explained by that component.
 - The first component explains approximately 13.3% of the variance.
 - The second component explains approximately 7.8% of the variance.
 - The third component explains approximately 5.9% of the variance.
 - And so on, with decreasing variance explained by each successive component.
- Cumulative explained variance:
 - This is the cumulative or total variance explained by adding up the individual components.
 - The cumulative explained variance helps us understand how much of the total variance is captured by including a certain number of components.
 - In this case, to retain 95% of the total variance, you would need to include 55 components.
 - The cumulative explained variance gradually increases as more components are added, reaching approximately 95% with the inclusion of these 55 components.

Question 2

After PCA, perform cluster analysis to identify cohorts, define these cohorts (cluster profiling), and specify the insights found.

Question 3

Highlight companies from different sectors falling into the same cohort, and share your findings.

```

from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

filtered_data = filtered_data.replace('[\$,]', '', regex=True)

```

```

# Standardize the data
numerical_data = filtered_data.drop(columns=['symbol', 'company',
'sector', 'subsector', 'Name', 'filename', 'Date'])
scaler = StandardScaler()
numerical_data_scaled = scaler.fit_transform(numerical_data)

# Determine the optimal number of clusters
inertia = []
K = range(2, 11)
for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans.fit(numerical_data_scaled)
    inertia.append(kmeans.inertia_)

# Plot the inertia
import matplotlib.pyplot as plt
plt.plot(K, inertia, 'bx-')
plt.xlabel('k')
plt.ylabel('Inertia')
plt.title('Inertia for different number of clusters')
plt.show()

# Choose the number of clusters based on the elbow method
optimal_k = int(input("Enter the optimal number of clusters based on
the plot: "))

# Perform k-means clustering with the optimal number of clusters
kmeans = KMeans(n_clusters=optimal_k, random_state=42)
clusters = kmeans.fit_predict(numerical_data_scaled)


# Add cluster labels to the original data
filtered_data['Cluster'] = clusters

# Cluster profiling
for i in range(optimal_k):
    cluster_data = filtered_data[filtered_data['Cluster'] == i]
    print(f"\nCluster {i}:")
    print(cluster_data.describe())

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value
of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value
of `n_init` explicitly to suppress the warning
    super()._check_params_vs_input(X, default_n_init=10)

```

```
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)  
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\sklearn\cluster\_kmeans.py:1416: FutureWarning: The default  
value of `n_init` will change from 10 to 'auto' in 1.4. Set the value  
of `n_init` explicitly to suppress the warning  
    super()._check_params_vs_input(X, default_n_init=10)
```



Enter the optimal number of clusters based on the plot: 5

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\cluster_kmeans.py:1416: FutureWarning: The default value of 'n_init' will change from 10 to 'auto' in 1.4. Set the value of 'n_init' explicitly to suppress the warning

super()._check_params_vs_input(X, default_n_init=10)

C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\1829115453.py:35: PerformanceWarning: DataFrame is highly fragmented. This is usually the result of calling 'frame.insert' many times, which has poor performance. Consider joining all columns at once using pd.concat(axis=1) instead. To get a de-fragmented frame, use 'newframe = frame.copy()'

filtered_data['Cluster'] = clusters

Cluster 0:

	asset_turnover_2017	asset_turnover_2018	asset_turnover_2019
count	72518.000000	72518.000000	72518.000000
mean	0.561683	0.537178	0.530919

min	0.220000	0.240000	0.220000
25%	0.340000	0.320000	0.310000
50%	0.420000	0.410000	0.400000
75%	0.660000	0.610000	0.650000
max	1.710000	1.960000	1.770000
std	0.347170	0.362933	0.343298

	asset_turnover_2020	asset_turnover_2021	asset_turnover_latest
\			
count	72518.000000	72518.000000	72518.000000
mean	0.481713	0.524100	0.139900
min	0.190000	0.200000	0.050000
25%	0.300000	0.320000	0.090000
50%	0.420000	0.420000	0.110000
75%	0.530000	0.710000	0.190000
max	1.580000	1.570000	0.350000
std	0.283483	0.293567	0.074177

	buyback_yield_2018	buyback_yield_2019	buyback_yield_2020	\
count	72518.000000	72518.000000	72518.000000	
mean	4.786190	3.341441	1.781781	
min	-24.970000	-2.690000	-3.190000	
25%	1.280000	1.230000	-0.020000	
50%	3.480000	2.484000	1.390000	
75%	7.790000	5.280000	2.220000	
max	25.030000	16.930000	14.680000	
std	8.707862	3.877649	3.337185	

	buyback_yield_2021	...	Net Change	Percentage Change	\
count	72518.000000	...	72518.000000	72518.000000	
mean	2.669988	...	1.540009	0.013463	
min	-1.340000	...	0.010000	-0.002500	
25%	0.040000	...	0.400000	0.004800	
50%	2.080000	...	0.820000	0.013600	
75%	3.056667	...	1.790000	0.018300	
max	17.620000	...	8.750000	0.047000	
std	4.096502	...	1.938724	0.011793	

		Date	High	Low \
count		72518	72518.000000	72518.000000
mean	2017-11-17 03:46:44.920157952		86.649985	84.806734
min	2012-10-31 00:00:00		2.710000	2.550000
25%	2015-05-29 00:00:00		36.131024	35.466250
50%	2017-11-20 00:00:00		63.630001	62.430000
75%	2020-05-18 00:00:00		101.040001	99.010002
max	2022-10-28 00:00:00		825.619995	817.530029
std		NaN	88.427590	86.356060

	Open	Close	Volume	Adj Close	Cluster
count	72518.000000	72518.000000	7.251800e+04	72518.000000	72518.0
mean	85.735961	85.747325	6.904498e+06	78.708468	0.0
min	2.640000	2.650000	0.000000e+00	2.363895	0.0
25%	35.813594	35.820000	2.087925e+06	30.612071	0.0
50%	63.029999	63.070000	3.844200e+06	55.068943	0.0
75%	100.070000	100.120003	7.891275e+06	90.639545	0.0
max	823.080017	821.010010	4.056765e+08	821.010010	0.0
std	87.395088	87.392881	9.560989e+06	86.655942	0.0

[8 rows x 222 columns]

Cluster 1:

	asset_turnover_2017	asset_turnover_2018
asset_turnover_2019 \		
count	90623.000000	90623.000000
mean	0.861884	0.905200
min	0.230000	0.300000
25%	0.640675	0.520000
50%	0.640675	0.700000
75%	1.016481	0.940000
max	3.710000	3.670000
std	0.602526	0.686774

	asset_turnover_2020	asset_turnover_2021	asset_turnover_latest	
\				
count	90623.000000	90623.000000	90623.000000	
mean	0.812194	0.816251	0.204677	
min	0.260000	0.200000	0.050000	
25%	0.480000	0.460000	0.120000	
50%	0.590000	0.630000	0.170000	
75%	0.830000	0.960000	0.250000	
max	3.300000	3.410000	0.830000	
std	0.614967	0.589761	0.140586	

	buyback_yield_2018	buyback_yield_2019	buyback_yield_2020	\
count	90623.000000	90623.000000	90623.000000	
mean	3.332830	2.635288	1.837107	
min	-6.700000	-6.560000	-5.110000	
25%	0.730000	0.190000	0.550000	
50%	2.940000	1.800000	1.430000	
75%	5.754573	4.140000	1.955939	
max	16.230000	13.910000	17.900000	
std	4.079243	3.987599	3.394102	

	buyback_yield_2021	...	Net Change	Percentage Change	\
count	90623.000000	...	90623.000000	90623.000000	
mean	1.485979	...	6.561785	0.026915	
min	-3.960000	...	0.130000	-0.015400	
25%	0.720000	...	1.660000	0.008500	
50%	1.510000	...	2.850000	0.025300	
75%	2.658531	...	7.315000	0.037800	
max	5.100000	...	45.730000	0.066800	
std	1.633559	...	9.906918	0.019669	

	Date	High	Low	\
count	90623	90623.000000	90623.000000	
mean	2018-02-15 17:23:46.962250240	176.377823	171.544548	
min	2012-10-31 00:00:00	1.690000	1.610000	
25%	2015-09-09 00:00:00	44.400002	43.247499	
50%	2018-04-24 00:00:00	89.099998	86.739998	
75%	2020-08-18 00:00:00	181.687500	176.235001	
max	2022-10-28 00:00:00	2715.659912	2632.219971	
std	NaN	303.617737	295.864383	

	Open	Close	Volume	Adj Close	Cluster
count	90623.000000	90623.000000	9.062300e+04	90623.000000	90623.0
mean	174.005086	173.995813	1.718028e+07	171.447794	1.0
min	1.620000	1.620000	0.000000e+00	1.620000	1.0
25%	43.799999	43.830002	1.276800e+06	40.087200	1.0
50%	87.900002	87.959999	2.677800e+06	84.903564	1.0
75%	178.980003	179.070000	1.081930e+07	175.985001	1.0
max	2680.000000	2703.260010	1.460852e+09	2703.260010	1.0
std	299.863290	299.711606	4.627886e+07	300.133791	0.0
[8 rows x 222 columns]					
Cluster 2:					
	asset_turnover_2017	asset_turnover_2018			
asset_turnover_2019 \					
count	69831.000000	69831.000000		69831.000000	
mean	0.724879	0.782842		0.751244	
min	0.140000	0.080000		0.030000	
25%	0.550000	0.550000		0.510000	
50%	0.650000	0.640000		0.640000	
75%	0.870000	0.920000		0.960000	
max	1.570000	1.590000		1.610000	
std	0.258528	0.312575		0.327476	
	asset_turnover_2020	asset_turnover_2021	asset_turnover_latest		
\count	69831.000000	69831.000000		69831.000000	
mean	0.687089	0.710768		0.175384	
min	0.060000	0.310000		0.070000	
25%	0.530000	0.490000		0.130000	

50%	0.620000	0.720000	0.150000
75%	0.770000	0.760000	0.200000
max	1.480000	1.460000	0.390000
std	0.281022	0.261902	0.083517

	buyback_yield_2018	buyback_yield_2019	buyback_yield_2020	\
count	69831.000000	69831.000000	69831.000000	
mean	1.422781	1.825195	1.001398	
min	-11.210000	-2.790000	-5.220000	
25%	0.610000	0.610000	0.190000	
50%	1.520000	1.770000	1.090000	
75%	2.897143	2.484000	1.651027	
max	5.290000	6.610000	7.650000	
std	2.737600	1.822117	2.263718	

	buyback_yield_2021	...	Net Change	Percentage Change	\
count	69831.000000	...	69831.000000	69831.000000	
mean	1.370484	...	4.207268	0.016215	
min	-1.810000	...	0.010000	-0.041700	
25%	0.680000	...	0.920000	0.005300	
50%	1.470000	...	2.350000	0.011000	
75%	1.958000	...	6.415000	0.027100	
max	4.830000	...	15.735000	0.063600	
std	1.149635	...	4.187162	0.021240	

		Date	High	Low	\
count		69831	69831.000000	69831.000000	
mean	2017-11-30 21:27:09.333677056		122.245129	118.849391	
min	2012-10-31 00:00:00		2.907500	2.787500	
25%	2015-05-28 00:00:00		43.500000	42.419998	
50%	2017-12-19 00:00:00		75.080002	73.470001	
75%	2020-06-10 12:00:00		157.404999	152.520004	
max	2022-10-28 00:00:00		761.049988	731.450012	
std		NaN	123.412805	119.736824	

	Open	Close	Volume	Adj Close	Cluster
count	69831.000000	69831.000000	6.983100e+04	69831.000000	69831.0
mean	120.585786	120.596341	9.074507e+06	119.478833	2.0
min	2.872500	2.845000	1.210000e+04	2.611672	2.0
25%	42.980000	42.980000	1.168600e+06	41.810579	2.0
50%	74.309998	74.300003	2.483200e+06	72.633209	2.0

75%	154.940002	154.939995	6.255800e+06	154.114937	2.0
max	752.559998	752.559998	4.643901e+08	752.559998	2.0
std	121.658382	121.616961	1.745319e+07	121.897363	0.0

[8 rows x 222 columns]

Cluster 3:

	asset_turnover_2017	asset_turnover_2018	
asset_turnover_2019 \			
count	2517.000000	2.517000e+03	2.517000e+03
mean	0.640675	4.800000e-01	4.600000e-01
min	0.640675	4.800000e-01	4.600000e-01
25%	0.640675	4.800000e-01	4.600000e-01
50%	0.640675	4.800000e-01	4.600000e-01
75%	0.640675	4.800000e-01	4.600000e-01
max	0.640675	4.800000e-01	4.600000e-01
std	0.000000	5.552218e-17	5.552218e-17

	asset_turnover_2020	asset_turnover_2021	asset_turnover_latest
\			
count	2517.00	2517.00	2517.00
mean	0.44	0.44	0.13
min	0.44	0.44	0.13
25%	0.44	0.44	0.13
50%	0.44	0.44	0.13
75%	0.44	0.44	0.13
max	0.44	0.44	0.13
std	0.00	0.00	0.00

	buyback_yield_2018	buyback_yield_2019	buyback_yield_2020	\
count	2.517000e+03	2517.0	2.517000e+03	
mean	1.400000e+00	1.5	4.860000e+00	

min	1.400000e+00	1.5	4.860000e+00
25%	1.400000e+00	1.5	4.860000e+00
50%	1.400000e+00	1.5	4.860000e+00
75%	1.400000e+00	1.5	4.860000e+00
max	1.400000e+00	1.5	4.860000e+00
std	4.441775e-16	0.0	8.883549e-16

	buyback_yield_2021	...	Net Change	Percentage Change \
count	2.517000e+03	...	2517.00	2.517000e+03
mean	3.030000e+00	...	3.75	2.320000e-02
min	3.030000e+00	...	3.75	2.320000e-02
25%	3.030000e+00	...	3.75	2.320000e-02
50%	3.030000e+00	...	3.75	2.320000e-02
75%	3.030000e+00	...	3.75	2.320000e-02
max	3.030000e+00	...	3.75	2.320000e-02
std	4.441775e-16	...	0.00	3.470136e-18

	Date	High	Low
Open \			
count	2517	2517.000000	2517.000000
2517.000000			
mean	2017-10-31 04:29:27.818831872	70.468668	68.385831
69.444189			
min	2012-10-31 00:00:00	13.553333	13.026667
13.200000			
25%	2015-05-04 00:00:00	42.076668	40.703335
41.500000			
50%	2017-10-30 00:00:00	57.419998	55.400002
56.606667			
75%	2020-05-01 00:00:00	80.936668	79.173332
80.073334			
max	2022-10-28 00:00:00	213.633331	208.103333
210.303329			
std	NaN	47.259005	45.673154
46.457123			

	Close	Volume	Adj Close	Cluster
count	2517.000000	2.517000e+03	2517.000000	2517.0
mean	69.447481	4.333416e+06	69.447481	3.0
min	13.186667	1.947000e+05	13.186667	3.0
25%	41.486668	2.718900e+06	41.486668	3.0
50%	56.536667	3.593100e+06	56.536667	3.0
75%	80.073334	4.847700e+06	80.073334	3.0
max	209.669998	6.535920e+07	209.669998	3.0
std	46.478426	3.313449e+06	46.478426	0.0

[8 rows x 222 columns]

Cluster 4:
asset_turnover_2017 asset_turnover_2018

asset_turnover_2019 \			
count	5.330000e+02	5.330000e+02	533.000000
mean	1.016481e+00	1.206477e+00	1.216292
min	1.016481e+00	1.206477e+00	1.216292
25%	1.016481e+00	1.206477e+00	1.216292
50%	1.016481e+00	1.206477e+00	1.216292
75%	1.016481e+00	1.206477e+00	1.216292
max	1.016481e+00	1.206477e+00	1.216292
std	2.222532e-16	2.222532e-16	0.000000
	asset_turnover_2020	asset_turnover_2021	asset_turnover_latest
\			
count	5.330000e+02	533.00	533.00
mean	1.063421e+00	0.01	0.01
min	1.063421e+00	0.01	0.01
25%	1.063421e+00	0.01	0.01
50%	1.063421e+00	0.01	0.01
75%	1.063421e+00	0.01	0.01
max	1.063421e+00	0.01	0.01
std	2.222532e-16	0.00	0.00
	buyback_yield_2018	buyback_yield_2019	buyback_yield_2020 \
count	533.000000	533.000000	533.000000
mean	5.754573	4.123958	1.955939
min	5.754573	4.123958	1.955939
25%	5.754573	4.123958	1.955939
50%	5.754573	4.123958	1.955939
75%	5.754573	4.123958	1.955939
max	5.754573	4.123958	1.955939
std	0.000000	0.000000	0.000000
	buyback_yield_2021	...	Net Change Percentage Change \
count	5.330000e+02	...	5.330000e+02 533.0000
mean	7.000000e-02	...	1.015000e+00 0.0793
min	7.000000e-02	...	1.015000e+00 0.0793

25%	7.000000e-02	...	1.015000e+00	0.0793
50%	7.000000e-02	...	1.015000e+00	0.0793
75%	7.000000e-02	...	1.015000e+00	0.0793
max	7.000000e-02	...	1.015000e+00	0.0793
std	1.389082e-17	...	2.222532e-16	0.0000

	Date	High	Low
Open \			
count	533	533.000000	533.000000
533.000000			
mean	2021-10-09 04:11:15.422138880	23.486144	21.645341
22.555747			
min	2020-09-18 00:00:00	9.640000	9.600000
9.630000			
25%	2021-03-31 00:00:00	17.240000	16.100000
16.500000			
50%	2021-10-08 00:00:00	22.100000	20.450001
21.340000			
75%	2022-04-20 00:00:00	27.139999	24.799999
25.850000			
max	2022-10-28 00:00:00	64.860001	56.080002
62.869999			
std	NaN	10.668913	9.221727
9.914020			

	Close	Volume	Adj Close	Cluster
count	533.000000	5.330000e+02	533.000000	533.0
mean	22.560932	2.854926e+07	22.560932	4.0
min	9.630000	4.450000e+04	9.630000	4.0
25%	16.600000	1.091830e+07	16.600000	4.0
50%	21.410000	1.863580e+07	21.410000	4.0
75%	25.549999	3.157320e+07	25.549999	4.0
max	58.049999	3.772209e+08	58.049999	4.0
std	9.969653	3.601109e+07	9.969653	0.0

[8 rows x 222 columns]

Cluster 0:

- Number of Data Points: 72,518
- Asset Turnover (2017-2021): Mean around 0.56 to 0.52
- Buyback Yield (2018-2021): Mean around 4.78 to 2.67
- Net Change: Mean around 1.54
- Percentage Change: Mean around 0.0135
- Price (High and Low): Mean High around 86.65, Mean Low around 84.81
- Volume: Mean around 6,904,498
- Adj Close: Mean around 78.71

Cluster 1:

- Number of Data Points: 90,623
- Asset Turnover (2017-2021): Mean around 0.86 to 0.81
- Buyback Yield (2018-2021): Mean around 3.33 to 1.48
- Net Change: Mean around 6.56
- Percentage Change: Mean around 0.0269
- Price (High and Low): Mean High around 176.38, Mean Low around 171.54
- Volume: Mean around 17,180,283
- Adj Close: Mean around 171.45

Cluster 2:

- Number of Data Points: 69,831
- Asset Turnover (2017-2021): Mean around 0.72 to 0.71
- Buyback Yield (2018-2021): Mean around 1.42 to 1.36
- Net Change: Mean around 4.21
- Percentage Change: Mean around 0.0162
- Price (High and Low): Mean High around 122.25, Mean Low around 118.85
- Volume: Mean around 9,074,507
- Adj Close: Mean around 119.48

Cluster 3:

- Number of Data Points: 2,517
- Asset Turnover (2017-2021): All values are approximately 0.64
- Buyback Yield (2018-2021): Buyback Yield is constant at 1.4 and 1.5
- Net Change: Constant at 3.75
- Percentage Change: Constant at 0.0232
- Price (High and Low): Mean High around 70.47, Mean Low around 68.39
- Volume: Mean around 4,333,416
- Adj Close: Mean around 69.45

Cluster 4:

- Number of Data Points: 533
- Asset Turnover (2017-2021): Asset Turnover is relatively stable around 1.02
- Buyback Yield (2018-2021): Buyback Yield is stable with small fluctuations around 5.75 and 4.12
- Net Change: Constant at 1.015
- Percentage Change: Constant at 0.0793
- Price (High and Low): Mean High around 23.49, Mean Low around 21.65
- Volume: Mean around 28,549,260
- Adj Close: Mean around 22.56

```
# Group by clusters and list unique sectors in each cluster
grouped = filtered_data.groupby('Cluster')['sector'].unique()

# Identify clusters with companies from multiple sectors
```

```

mixed_clusters = {k: v for k, v in grouped.items() if len(v) > 1}

# Print the findings
print("Clusters with companies from multiple sectors:")
for cluster, sectors in mixed_clusters.items():
    #print(f"Cluster {cluster} has companies from sectors: {'',
    '.join(sectors)}")
    pass
# Highlight companies from different sectors in the same cluster
for cluster, sectors in mixed_clusters.items():
    cluster_data = filtered_data[filtered_data['Cluster'] == cluster]
    #print(f"\nCompanies in Cluster {cluster}:")
    for sector in sectors:
        sector_companies = cluster_data[cluster_data['sector'] ==
sector]['company'].tolist()
        print(sector_companies[0:5])
        #print(f"Sector {sector}: {'', '.join(sector_companies)}")

```

```

Clusters with companies from multiple sectors:
['American Electric Power', 'American Electric Power', 'American
Electric Power', 'American Electric Power', 'American Electric Power']
['Amgen', 'Amgen', 'Amgen', 'Amgen', 'Amgen']
['Broadcom Inc.', 'Broadcom Inc.', 'Broadcom Inc.', 'Broadcom Inc.',
'Broadcom Inc.']
['Charter Communications', 'Charter Communications', 'Charter
Communications', 'Charter Communications', 'Charter Communications']
['CSX Corporation', 'CSX Corporation', 'CSX Corporation', 'CSX
Corporation', 'CSX Corporation']
['eBay', 'eBay', 'eBay', 'eBay', 'eBay']
['Keurig Dr Pepper', 'Keurig Dr Pepper', 'Keurig Dr Pepper', 'Keurig
Dr Pepper', 'Keurig Dr Pepper']
['Apple Inc.', 'Apple Inc.', 'Apple Inc.', 'Apple Inc.', 'Apple Inc.']
['Airbnb', 'Airbnb', 'Airbnb', 'Airbnb', 'Airbnb']
['Baidu', 'Baidu', 'Baidu', 'Baidu', 'Baidu']
['Biogen', 'Biogen', 'Biogen', 'Biogen', 'Biogen']
['Costco', 'Costco', 'Costco', 'Costco', 'Costco']
['Cintas', 'Cintas', 'Cintas', 'Cintas', 'Cintas']
['Adobe Inc.', 'Adobe Inc.', 'Adobe Inc.', 'Adobe Inc.', 'Adobe Inc.']
['Align Technology', 'Align Technology', 'Align Technology', 'Align
Technology', 'Align Technology']
['Activision Blizzard', 'Activision Blizzard', 'Activision Blizzard',
'Activision Blizzard', 'Activision Blizzard']
['Copart', 'Copart', 'Copart', 'Copart', 'Copart']
['Lululemon', 'Lululemon', 'Lululemon', 'Lululemon', 'Lululemon']
['Monster Beverage', 'Monster Beverage', 'Monster Beverage', 'Monster
Beverage', 'Monster Beverage']

```

Question 4

Plot seasonality, trend, and irregular components over time for the historical stock price of Apple.

```
# Filter data for AAPL
aapl_data = filtered_data[filtered_data['symbol'] == 'AAPL']

# Convert 'Date' to datetime and set as index
aapl_data['Date'] = pd.to_datetime(aapl_data['Date'])
aapl_data.set_index('Date', inplace=True)

# Sort the data by date
aapl_data.sort_index(inplace=True)

# Perform seasonal decomposition
result = seasonal_decompose(aapl_data['Close'],
                             model='multiplicative', period=252) # assuming daily data with annual seasonality

# Plot the decomposition
fig, (ax1, ax2, ax3, ax4) = plt.subplots(4, 1, figsize=(15, 8))

result.observed.plot(ax=ax1)
ax1.set_ylabel('Observed')

result.trend.plot(ax=ax2)
ax2.set_ylabel('Trend')


result.seasonal.plot(ax=ax3)
ax3.set_ylabel('Seasonal')

result.resid.plot(ax=ax4)
ax4.set_ylabel('Residual')

plt.tight_layout()
plt.show()
```

C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\3082912243.py:5:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
aapl_data['Date'] = pd.to_datetime(aapl_data['Date'])



Question 5

Based on trend and seasonality, choose an appropriate exponential smoothing method to forecast the weekend share price value for the next 12 months.

```
# Filter data for AAPL and close prices
apl_data = filtered_data[filtered_data['symbol'] == 'AAPL']

# Convert 'Date' to datetime and set as index
apl_data['Date'] = pd.to_datetime(apl_data['Date'])
apl_data.set_index('Date', inplace=True)

# Sort the data by date
apl_data.sort_index(inplace=True)

# Extract the 'Close' price series
close_prices = apl_data['Close']

# Fit the Holt-Winters exponential smoothing model
model = ExponentialSmoothing(close_prices, seasonal='multiplicative',
seasonal_periods=252)
fit = model.fit()

# Forecast the next 12 months (assuming 252 trading days in a year)
forecast = fit.forecast(steps=12)

# Create a new dataframe for forecasting with future dates
forecast_index = pd.date_range(start=close_prices.index[-1],
```

```
periods=13, freq='B') # Forecasting 12 months, plus one extra
forecast_df = pd.DataFrame({'Forecasted': forecast},
index=forecast_index)
```

```
# Concatenate the observed and forecasted data
combined_data = pd.concat([close_prices, forecast_df])
```

```
# Plot the observed and forecasted values
```

```
plt.figure(figsize=(10, 5))
plt.plot(combined_data.index, combined_data, label='Observed',
color='blue')
plt.plot(forecast_df.index, forecast_df['Forecasted'],
label='Forecasted', color='red', linestyle='--')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.title('Holt-Winters Forecast for AAPL Close Prices')
plt.legend()
plt.show()
```

C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\3877550715.py:5:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
aapl_data['Date'] = pd.to_datetime(aapl_data['Date'])
```

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
self._init_dates(dates, freq)
```

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\tsa\holtwinters\model.py:917: ConvergenceWarning: Optimization failed to converge. Check mle_retvals.


```
warnings.warn(
```

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
return get_prediction_index(
```

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:836: FutureWarning: No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
return get_prediction_index(
```



Question 6

Perform an augmented Dickey–Fuller test (ADF) to check for the stationarity of Apple stock.

```
# Perform Augmented Dickey-Fuller test
result = adfuller(close_prices)
print('ADF Statistic: %f' % result[0])
print('p-value: %f' % result[1])
print('Critical Values:')
for key, value in result[4].items():
    print('\t%s: %.3f' % (key, value))

# Interpretation
if result[1] <= 0.05:
    print("Reject the null hypothesis. The time series is stationary.")
else:
    print("Fail to reject the null hypothesis. The time series is not stationary.")

ADF Statistic: 0.294860
p-value: 0.977107
Critical Values:
    1%: -3.433
    5%: -2.863
   10%: -2.567
Fail to reject the null hypothesis. The time series is not stationary.
```

Insights & Summary:

- **ADF Statistic:** 0.294860
- **p-value:** 0.977107
- **Critical Values:**
 - 1%: -3.433
 - 5%: -2.863
 - 10%: -2.567

The Augmented Dickey-Fuller (ADF) test is commonly used to assess the stationarity of a time series. In this case, the ADF statistic is 0.294860, and the associated p-value is 0.977107.

The ADF statistic is used to compare against critical values at various significance levels (1%, 5%, and 10%) to determine whether the time series is stationary or not. In this instance, the critical values are as follows:

- 1%: -3.433
- 5%: -2.863
- 10%: -2.567

The null hypothesis of the ADF test is that the time series is non-stationary (it has a unit root). The test results indicate that we fail to reject the null hypothesis. In other words, there is insufficient evidence to conclude that the time series is stationary.

In practical terms, this suggests that the time series data exhibits some form of trend or seasonality, and it may require differencing or other transformations to achieve stationarity for certain statistical analyses or modeling purposes.

Question 7

Analyze the ACF and PACF plots for Apple's historical stock prices, strategize for ARIMA modeling, determine the appropriate values of p, d, and q, and forecast the month-end share price value for the next 12 months.

```
# Assuming filtered_data is already loaded and contains the data
apple_data = filtered_data[filtered_data['symbol'] == 'AAPL']
close_prices = apple_data['Close']

# Plot ACF and PACF
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
plot_acf(close_prices, ax=ax1)
plot_pacf(close_prices, ax=ax2)
plt.show()

# After analyzing the ACF and PACF plots, you can determine the values
for p, d, and q.
# For this example, I'll use placeholder values, but you should
replace them based on your observations.
```



```


p, d, q = 1, 1, 1 # Replace these values based on ACF and PACF plots

# Fit ARIMA model
model = ARIMA(close_prices, order=(p, d, q))
model_fit = model.fit()
print(model_fit.summary())

# Make predictions
forecast = model_fit.get_forecast(steps=12)
forecast_index = pd.date_range(start=apple_data['Date'].iloc[-1],
periods=13)[1:] # Exclude the first date
forecast_values = forecast.predicted_mean
forecast_conf_int = forecast.conf_int()

# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(apple_data['Date'], close_prices, label='Observed')
plt.plot(forecast_index, forecast_values, label='Forecast',
color='red')
plt.fill_between(forecast_index, forecast_conf_int.iloc[:, 0],
forecast_conf_int.iloc[:, 1], color='red', alpha=0.3)
plt.xlabel('Time')
plt.ylabel('Close Price')
plt.title('Apple Stock Price Forecast')
plt.legend()
plt.show()

```



C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: An unsupported index was provided and will be ignored when e.g. forecasting.

```
self._init_dates(dates, freq)
```

C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: An unsupported index was provided and will be ignored when e.g.

```

forecasting.
    self._init_dates(dates, freq)
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-
packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: An
unsupported index was provided and will be ignored when e.g.
forecasting.
    self._init_dates(dates, freq)

```

SARIMAX Results

```

=====
=====
Dep. Variable:                Close    No. Observations:
2517
Model:                        ARIMA(1, 1, 1)    Log Likelihood    -
4642.569
Date:                        Mon, 23 Oct 2023    AIC
9291.138
Time:                        18:07:38    BIC
9308.630
Sample:                        0    HQIC
9297.487
                                - 2517

Covariance Type:                opg

=====
=====
                                coef    std err          z      P>|z|      [0.025
0.975]
-----
ar.L1                0.2879        0.129        2.231        0.026        0.035
0.541
ma.L1               -0.3544        0.127       -2.793        0.005       -0.603
-0.106
sigma2                2.3456        0.028       84.351        0.000        2.291
2.400
=====
=====
Ljung-Box (L1) (Q):                0.00    Jarque-Bera (JB):
9669.23
Prob(Q):                0.99    Prob(JB):
0.00
Heteroskedasticity (H):                41.90    Skew:
-0.12
Prob(H) (two-sided):                0.00    Kurtosis:
12.60
=====
=====

```

Warnings:


```
[1] Covariance matrix calculated using the outer product of gradients  
(complex-step).
```

```
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\statsmodels\tsa\base\tsa_model.py:836: ValueWarning: No  
supported index is available. Prediction results will be given with an  
integer index beginning at `start`.
```

```
    return get_prediction_index(
```

```
C:\Users\bhara\AppData\Local\Programs\Python\Python311\Lib\site-  
packages\statsmodels\tsa\base\tsa_model.py:836: FutureWarning: No  
supported index is available. In the next version, calling this method  
in a model without a supported index will result in an exception.
```

```
    return get_prediction_index(
```



SARIMAX Results Detailed Summary

Model Configuration

- The SARIMAX (Seasonal AutoRegressive Integrated Moving Average with eXogenous factors) analysis was conducted using a specific model configuration: ARIMA(1, 1, 1).

- The analysis focused on the "Close" variable and utilized a dataset comprising 2517 observations.
- The log-likelihood of this model was computed to be -4642.569, providing a measure of how well the model fits the data.
- Two commonly used information criteria were computed: the Akaike Information Criterion (AIC) was found to be 9291.138, and the Bayesian Information Criterion (BIC) was 9308.630. These criteria help in model selection, considering the trade-off between goodness of fit and model complexity.
- The analysis was conducted on Mon, 23 Oct 2023, at 11:28:17. The covariance type used for this analysis was "opg" (Outer Product of Gradients).

Parameter Coefficients

- The model estimated the autoregressive coefficient (ar.L1) to be approximately 0.2879. This coefficient indicates the strength and direction of the linear relationship between the current observation and the previous observation in the time series.
- The standard error associated with ar.L1 was 0.129, reflecting the precision of this estimate.
- The z-statistic for ar.L1 was 2.231, and the corresponding p-value was 0.026. This p-value suggests that ar.L1 is statistically significant at the 0.05 significance level, indicating that the lag-1 autocorrelation is not zero.
- The 95% confidence interval for ar.L1 was calculated as [0.035, 0.541], providing a range of plausible values for this coefficient.
- Similarly, the model estimated the moving average coefficient (ma.L1) to be approximately -0.3544, with a standard error of 0.127. The associated z-statistic was -2.793, and the p-value was 0.005, indicating statistical significance.
- The 95% confidence interval for ma.L1 was [-0.603, -0.106].
- The model estimated the variance of the residuals (sigma2) to be approximately 2.3456, with a standard error of 0.028. The extremely high z-statistic of 84.351 and the very low p-value of 0.000 indicate that this estimate is highly statistically significant.
- The 95% confidence interval for sigma2 was [2.291, 2.400], providing a range for the variance estimate.

Diagnostic Tests

- The Ljung-Box test conducted at lag 1 (L1) resulted in a Q-statistic of 0.00, with a probability (Prob(Q)) of 0.99. This indicates no significant autocorrelation in the residuals at lag 1, suggesting that the model captures the temporal dependencies effectively.
- The Jarque-Bera (JB) statistic, which assesses the normality of the residuals, was found to be 9669.23, with a probability (Prob(JB)) of 0.00. The extremely low p-value suggests that the residuals do not follow a normal distribution.
- The heteroskedasticity test produced a high statistic of 41.90, and the two-sided probability was 0.00. This indicates that the variance of the residuals is not constant, suggesting the presence of heteroskedasticity.

Additional Information

- The skewness (Skew) of the residuals was computed to be -0.12, indicating a slight leftward skew in the distribution of residuals.
- The kurtosis (Kurtosis) of the residuals was found to be 12.60, suggesting heavy tails in the distribution. This implies that the residuals exhibit more extreme values than a normal distribution would predict.

Warnings

- It's important to note that the covariance matrix used in the analysis was calculated using the outer product of gradients (complex-step). This method may have implications for the precision of parameter estimates.

Interpretation

The SARIMAX model with an ARIMA(1, 1, 1) configuration was applied to the "Close" time series data. The results indicate that there is statistically significant autocorrelation at lag 1, as evidenced by the ar.L1 and ma.L1 coefficients. However, diagnostic tests reveal that the residuals do not follow a normal distribution and exhibit heteroskedasticity. Further refinement of the model or consideration of alternative models may be necessary to account for these issues and improve model performance. Additionally, the high kurtosis suggests that extreme values may occur more frequently than expected, which should be taken into account when interpreting the model's predictions.

Question 8

Find the mean absolute percentage error (MAPE) for a 12-month period to validate the model.

```
# Calculate MAPE
actual_values = close_prices[-12:]
mape = np.mean(np.abs((actual_values.values - forecast_values.values)
/ actual_values.values)) * 100
print(f'MAPE: {mape}')
```

MAPE: 6.198755830729437

Mean Absolute Percentage Error (MAPE): 6.20%

The Mean Absolute Percentage Error (MAPE) is a metric that measures the average percentage difference between observed and predicted values. In this context, a MAPE of 6.20% indicates that, on average, the model's predictions deviate from the actual values by about 6.20%. Lower MAPE values signify better predictive accuracy, while higher values indicate larger prediction errors. In this case, a MAPE of 6.20% suggests a reasonably good level of predictive accuracy.

```
filtered_data.columns
```

```
Index(['symbol', 'company', 'sector', 'subsector',
'asset_turnover_2017',
```

```

        'asset_turnover_2018', 'asset_turnover_2019',
'asset_turnover_2020',
        'asset_turnover_2021', 'asset_turnover_latest',
        '...',
        'Percentage Change', 'Date', 'High', 'Low', 'Open', 'Close',
'Volume',
        'Adj Close', 'filename', 'Cluster'],
        dtype='object', length=228)

```

Question 9

Identify the top 2 companies from each sector based on market capitalization, create trend charts for the month-end share price for the last five years (using the variable "adjusted close"), display the 12-month rolling mean and standard deviation in the same chart, and share your observations regarding the stationarity of all companies.

```

# Find top 2 companies from each sector based on market capitalization
top_companies = filtered_data.groupby('sector')[['company', 'Market
Cap']].apply(lambda x: x.nlargest(2, 'Market Cap'))
top_companies

```

sector		company	Market Cap
Communication Services	207656	Alphabet Inc. (Class C)	1355597700000
	207657	Alphabet Inc. (Class C)	1355597700000
Consumer Discretionary	60881	Amazon	1228246601480
	60882	Amazon	1228246601480
Consumer Staples	380803	PepsiCo	245342455939
	380804	PepsiCo	245342455939
Health Care	87193	AstraZeneca	173545001728
	87194	AstraZeneca	173545001728
Industrials	222565	Honeywell	126883210984
	222566	Honeywell	126883210984
Information Technology	3228	Apple Inc.	2625740143000
	3229	Apple Inc.	2625740143000
Utilities	32428	American Electric Power	44807878084
	32429	American Electric Power	44807878084

```

def stationarity_check(timeseries):
    # Perform Augmented Dickey-Fuller test
    print('Results of Dickey-Fuller Test:')
    dfctest = adfuller(timeseries, autolag='AIC')
    dfctest = pd.Series(dfctest[0:4], index=['Test Statistic', 'p-
value', '#Lags Used', 'Number of Observations Used'])
    for key, value in dfctest[4].items():
        dfctest['Critical Value (%s)' % key] = value
    print(dfctest)
    if dfctest['p-value'] <= 0.05:

```

```

        print("The series is stationary")
    else:
        print("The series is not stationary")

# Plot trend charts for the top 2 companies
for sector, companies in top_companies.groupby('sector'):
    print(f'Sector: {sector}')
    for company in companies['company'].unique():
        company_data = filtered_data[(filtered_data['sector'] ==
sector) & (filtered_data['company'] == company)]
        company_data = company_data.set_index('Date')
        company_data = company_data.sort_index()
        company_data_last_five_years = company_data.last('5Y')

        # Calculate 12-month rolling mean and standard deviation
        rolling_mean = company_data_last_five_years['Adj
Close'].rolling(window=12).mean()
        rolling_std = company_data_last_five_years['Adj
Close'].rolling(window=12).std()


        # Plot trend chart
        plt.figure(figsize=(10, 6))
        plt.plot(company_data_last_five_years['Adj Close'],
label='Month-end Share Price')
        plt.plot(rolling_mean, label='12-month Rolling Mean',
color='red')
        plt.plot(rolling_std, label='12-month Rolling Std',
color='green')
        plt.xlabel('Date')
        plt.ylabel('Adjusted Close Price')
        plt.title(f'Trend Chart for {company} ({sector})')
        plt.legend()
        plt.show()

        # Check stationarity
        stationarity_check(company_data_last_five_years['Adj Close'])

```

Sector: Communication Services

C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:
FutureWarning: last is deprecated and will be removed in a future
version. Please create a mask and filter using `.loc` instead
 company_data_last_five_years = company_data.last('5Y')



Results of Dickey-Fuller Test:

Test Statistic	-1.019369
p-value	0.746135
#Lags Used	9.000000
Number of Observations Used	1206.000000
Critical Value (1%)	-3.435784
Critical Value (5%)	-2.863940
Critical Value (10%)	-2.568048

dtype: float64

The series is not stationary

Sector: Consumer Discretionary

```
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:
FutureWarning: last is deprecated and will be removed in a future
version. Please create a mask and filter using `.loc` instead
  company_data_last_five_years = company_data.last('5Y')
```




Results of Dickey-Fuller Test:


Test Statistic	-1.840332
p-value	0.360630
#Lags Used	0.000000
Number of Observations Used	1215.000000
Critical Value (1%)	-3.435744
Critical Value (5%)	-2.863922
Critical Value (10%)	-2.568038

dtype: float64

The series is not stationary

Sector: Consumer Staples

```
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:
FutureWarning: last is deprecated and will be removed in a future
version. Please create a mask and filter using `.loc` instead
  company_data_last_five_years = company_data.last('5Y')
```



Results of Dickey-Fuller Test:


Test Statistic	-0.370841
p-value	0.914855
#Lags Used	9.000000
Number of Observations Used	1206.000000
Critical Value (1%)	-3.435784
Critical Value (5%)	-2.863940
Critical Value (10%)	-2.568048

dtype: float64

The series is not stationary

Sector: Health Care

C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:
 FutureWarning: last is deprecated and will be removed in a future
 version. Please create a mask and filter using `.loc` instead
 company_data_last_five_years = company_data.last('5Y')




Results of Dickey-Fuller Test:

Test Statistic	-1.355530
p-value	0.603401
#Lags Used	10.000000
Number of Observations Used	1205.000000
Critical Value (1%)	-3.435788
Critical Value (5%)	-2.863942
Critical Value (10%)	-2.568049

dtype: float64
The series is not stationary
Sector: Industrials

```
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:
FutureWarning: last is deprecated and will be removed in a future
version. Please create a mask and filter using `.loc` instead
    company_data_last_five_years = company_data.last('5Y')
```



Results of Dickey-Fuller Test:


Test Statistic	-1.407996
p-value	0.578437
#Lags Used	8.000000
Number of Observations Used	1207.000000
Critical Value (1%)	-3.435779
Critical Value (5%)	-2.863938
Critical Value (10%)	-2.568046

dtype: float64

The series is not stationary

Sector: Information Technology

```
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:  
FutureWarning: last is deprecated and will be removed in a future  
version. Please create a mask and filter using `.loc` instead  
    company_data_last_five_years = company_data.last('5Y')
```




Results of Dickey-Fuller Test:

Test Statistic	-0.734874
p-value	0.837509
#Lags Used	22.000000
Number of Observations Used	1193.000000
Critical Value (1%)	-3.435843
Critical Value (5%)	-2.863966
Critical Value (10%)	-2.568061

dtype: float64
The series is not stationary
Sector: Utilities

```
C:\Users\bhara\AppData\Local\Temp\ipykernel_1476\491218034.py:21:  
FutureWarning: last is deprecated and will be removed in a future  
version. Please create a mask and filter using `.loc` instead  
    company_data_last_five_years = company_data.last('5Y')
```



```
Results of Dickey-Fuller Test:
Test Statistic          -1.965287
p-value                  0.301998
#Lags Used               13.000000
Number of Observations Used 1202.000000
Critical Value (1%)      -3.435802
Critical Value (5%)      -2.863948
Critical Value (10%)     -2.568052
dtype: float64
The series is not stationary
```

Sector: Communication Services

- Data for the Communication Services sector was analyzed.
- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.
- The test resulted in a Test Statistic of -1.019369 and a p-value of 0.746135.
- A total of 9 lags were used in the test, and there were 1206 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435784, -2.863940, and -2.568048, respectively.
- Based on the results, the series is found to be non-stationary.

Sector: Consumer Discretionary

- Data for the Consumer Discretionary sector was analyzed.
- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.

- The test resulted in a Test Statistic of -1.840332 and a p-value of 0.360630.
- No lags were used in this test, and there were 1215 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435744, -2.863922, and -2.568038, respectively.
- Based on the results, the series is found to be non-stationary.

Sector: Consumer Staples

- Data for the Consumer Staples sector was analyzed.
- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.
- The test resulted in a Test Statistic of -0.370841 and a p-value of 0.914855.
- A total of 9 lags were used in the test, and there were 1206 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435784, -2.863940, and -2.568048, respectively.
- Based on the results, the series is found to be non-stationary.

Sector: Health Care

- Data for the Health Care sector was analyzed.
- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.
- The test resulted in a Test Statistic of -1.355530 and a p-value of 0.603401.
- A total of 10 lags were used in the test, and there were 1205 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435788, -2.863942, and -2.568049, respectively.
- Based on the results, the series is found to be non-stationary.

Sector: Industrials

- Data for the Industrials sector was analyzed.
- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.
- The test resulted in a Test Statistic of -1.407996 and a p-value of 0.578437.
- A total of 8 lags were used in the test, and there were 1207 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435779, -2.863938, and -2.568046, respectively.
- Based on the results, the series is found to be non-stationary.

Sector: Information Technology

- Data for the Information Technology sector was analyzed.
- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.
- The test resulted in a Test Statistic of -0.734874 and a p-value of 0.837509.
- A total of 22 lags were used in the test, and there were 1193 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435843, -2.863966, and -2.568061, respectively.
- Based on the results, the series is found to be non-stationary.

Sector: Utilities

- Data for the Utilities sector was analyzed.

- The Dickey-Fuller Test was conducted to assess the stationarity of the time series data.
- The test resulted in a Test Statistic of -1.965287 and a p-value of 0.301998.
- A total of 13 lags were used in the test, and there were 1202 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.435802, -2.863948, and -2.568052, respectively.
- Based on the results, the series is found to be non-stationary.

Question 10 :

Conduct an ADF test to verify the stationarity of the companies selected in the previous step

```
# Define a function to perform the ADF test
def adf_test(series, title=''):
    print(f'Augmented Dickey-Fuller Test: {title}')
    result = adfuller(series, autolag='AIC')
    labels = ['ADF Test Statistic', 'p-value', '# Lags Used', '#
Observations Used']
    out = pd.Series(result[0:4], index=labels)

    for key, val in result[4].items():
        out[f'Critical Value ({key})'] = val

    print(out.to_string())
    if result[1] <= 0.05:
        print("The data is stationary")
    else:
        print("The data is not stationary")

# Replace 'dataframe' with your pandas DataFrame containing the
companies' stock prices
for company in top_companies['company'].values:
    series = filtered_data[filtered_data['company'] == company]['Adj
Close']
    adf_test(series, title=company)
```

```
Augmented Dickey-Fuller Test: Alphabet Inc. (Class C)
ADF Test Statistic      -0.772851
p-value                  0.826967
# Lags Used              24.000000
# Observations Used     2492.000000
Critical Value (1%)     -3.432977
Critical Value (5%)     -2.862701
Critical Value (10%)    -2.567388
The data is not stationary
Augmented Dickey-Fuller Test: Alphabet Inc. (Class C)
ADF Test Statistic      -0.772851
p-value                  0.826967
# Lags Used              24.000000
# Observations Used     2492.000000
```



```

Critical Value (1%)      -3.432977
Critical Value (5%)      -2.862701
Critical Value (10%)     -2.567388
The data is not stationary
Augmented Dickey-Fuller Test: Amazon
ADF Test Statistic      -1.046041
p-value                  0.736104
# Lags Used              27.000000
# Observations Used      2489.000000
Critical Value (1%)      -3.432980
Critical Value (5%)      -2.862702
Critical Value (10%)     -2.567389
The data is not stationary
Augmented Dickey-Fuller Test: Amazon
ADF Test Statistic      -1.046041
p-value                  0.736104
# Lags Used              27.000000
# Observations Used      2489.000000
Critical Value (1%)      -3.432980
Critical Value (5%)      -2.862702
Critical Value (10%)     -2.567389
The data is not stationary
Augmented Dickey-Fuller Test: PepsiCo
ADF Test Statistic      0.445183
p-value                  0.983108
# Lags Used              27.000000
# Observations Used      2489.000000
Critical Value (1%)      -3.432980
Critical Value (5%)      -2.862702
Critical Value (10%)     -2.567389
The data is not stationary
Augmented Dickey-Fuller Test: PepsiCo
ADF Test Statistic      0.445183
p-value                  0.983108
# Lags Used              27.000000
# Observations Used      2489.000000
Critical Value (1%)      -3.432980
Critical Value (5%)      -2.862702
Critical Value (10%)     -2.567389
The data is not stationary
Augmented Dickey-Fuller Test: AstraZeneca
ADF Test Statistic      -0.734786
p-value                  0.837532
# Lags Used              10.000000
# Observations Used      2506.000000
Critical Value (1%)      -3.432962
Critical Value (5%)      -2.862694
Critical Value (10%)     -2.567384
The data is not stationary

```

Augmented Dickey-Fuller Test: AstraZeneca

ADF Test Statistic -0.734786

p-value 0.837532

Lags Used 10.000000

Observations Used 2506.000000

Critical Value (1%) -3.432962

Critical Value (5%) -2.862694

Critical Value (10%) -2.567384

The data is not stationary

Augmented Dickey-Fuller Test: Honeywell

ADF Test Statistic -0.987240

p-value 0.757890

Lags Used 9.000000

Observations Used 2507.000000

Critical Value (1%) -3.432961

Critical Value (5%) -2.862694

Critical Value (10%) -2.567384

The data is not stationary

Augmented Dickey-Fuller Test: Honeywell

ADF Test Statistic -0.987240

p-value 0.757890

Lags Used 9.000000

Observations Used 2507.000000

Critical Value (1%) -3.432961

Critical Value (5%) -2.862694

Critical Value (10%) -2.567384

The data is not stationary

Augmented Dickey-Fuller Test: Apple Inc.

ADF Test Statistic 0.345148

p-value 0.979324

Lags Used 22.000000

Observations Used 2494.000000

Critical Value (1%) -3.432975

Critical Value (5%) -2.862700

Critical Value (10%) -2.567387

The data is not stationary

Augmented Dickey-Fuller Test: Apple Inc.

ADF Test Statistic 0.345148

p-value 0.979324

Lags Used 22.000000

Observations Used 2494.000000

Critical Value (1%) -3.432975

Critical Value (5%) -2.862700

Critical Value (10%) -2.567387

The data is not stationary

Augmented Dickey-Fuller Test: American Electric Power

ADF Test Statistic -1.180378

p-value 0.681979

Lags Used 13.000000

```

# Observations Used      2503.000000
Critical Value (1%)      -3.432965
Critical Value (5%)      -2.862695
Critical Value (10%)     -2.567385
The data is not stationary
Augmented Dickey-Fuller Test: American Electric Power
ADF Test Statistic       -1.180378
p-value                  0.681979
# Lags Used              13.000000
# Observations Used      2503.000000
Critical Value (1%)      -3.432965
Critical Value (5%)      -2.862695
Critical Value (10%)     -2.567385
The data is not stationary

```

Augmented Dickey-Fuller Test: Alphabet Inc. (Class C)

- The Augmented Dickey-Fuller (ADF) Test was conducted on Alphabet Inc. (Class C) stock data.
- The ADF Test Statistic was found to be -0.772851, and the corresponding p-value was 0.826967.
- 24 lags were used in the test, and there were 2492 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432977, -2.862701, and -2.567388, respectively.
- Based on the results, the data is not stationary.

Augmented Dickey-Fuller Test: Amazon

- The Augmented Dickey-Fuller (ADF) Test was conducted on Amazon stock data.
- The ADF Test Statistic was found to be -1.046041, and the corresponding p-value was 0.736104.
- 27 lags were used in the test, and there were 2489 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432980, -2.862702, and -2.567389, respectively.
- Based on the results, the data is not stationary.

Augmented Dickey-Fuller Test: PepsiCo

- The Augmented Dickey-Fuller (ADF) Test was conducted on PepsiCo stock data.
- The ADF Test Statistic was found to be 0.445183, and the corresponding p-value was 0.983108.
- 27 lags were used in the test, and there were 2489 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432980, -2.862702, and -2.567389, respectively.
- Based on the results, the data is not stationary.

Augmented Dickey-Fuller Test: AstraZeneca

- The Augmented Dickey-Fuller (ADF) Test was conducted on AstraZeneca stock data.

- The ADF Test Statistic was found to be -0.734786, and the corresponding p-value was 0.837532.
- 10 lags were used in the test, and there were 2506 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432962, -2.862694, and -2.567384, respectively.
- Based on the results, the data is not stationary.

Augmented Dickey-Fuller Test: Honeywell

- The Augmented Dickey-Fuller (ADF) Test was conducted on Honeywell stock data.
- The ADF Test Statistic was found to be -0.987240, and the corresponding p-value was 0.757890.
- 9 lags were used in the test, and there were 2507 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432961, -2.862694, and -2.567384, respectively.
- Based on the results, the data is not stationary.

Augmented Dickey-Fuller Test: Apple Inc.

- The Augmented Dickey-Fuller (ADF) Test was conducted on Apple Inc. stock data.
- The ADF Test Statistic was found to be 0.345148, and the corresponding p-value was 0.979324.
- 22 lags were used in the test, and there were 2494 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432975, -2.862700, and -2.567387, respectively.
- Based on the results, the data is not stationary.

Augmented Dickey-Fuller Test: American Electric Power

- The Augmented Dickey-Fuller (ADF) Test was conducted on American Electric Power stock data.
- The ADF Test Statistic was found to be -1.180378, and the corresponding p-value was 0.681979.
- 13 lags were used in the test, and there were 2503 observations.
- Critical values for significance levels of 1%, 5%, and 10% were -3.432965, -2.862695, and -2.567385, respectively.
- Based on the results, the data is not stationary.

```
filtered_data.to_csv('filtered_data.csv')
```

Question 11

Perform batch forecasting for the top 2 companies from each sector based on market capitalization for the weekend share price value for the next 12 months using Auto ARIMA, and find the MAPE for a 12-month period to validate the model.

```
import warnings
warnings.filterwarnings('ignore')
```

```

from sklearn.metrics import mean_absolute_percentage_error

# Initialize a DataFrame to store MAPE results
mape_results = []

# Define a function to calculate MAPE
def calculate_mape(actual, forecast):
    return mean_absolute_percentage_error(actual, forecast)

# Group companies by sector and company
sector_grouped = filtered_data.groupby(['sector', 'company'])

# Calculate the total Market Cap for each company in each sector
company_market_cap = sector_grouped['Market Cap'].sum().reset_index()

# Sort companies by Market Cap within each sector
company_market_cap_sorted =
company_market_cap.sort_values(by=['sector', 'Market Cap'],
ascending=[True, False])

# Select the top 2 companies in each sector
top_companies = company_market_cap_sorted.groupby('sector').head(2)

# Iterate through the top companies in each sector
for index, row in top_companies.iterrows():
    sector = row['sector']
    company_name = row['company']

    # Filter data for the selected company and sector
    company_data = filtered_data[(filtered_data['company'] ==
company_name) & (filtered_data['sector'] == sector)]

    # Extract the share price data for modeling
    share_price = company_data[['Date', 'Close']]
    share_price.set_index('Date', inplace=True)

    # Perform train-test split
    train_size = len(share_price) - 12 # Use all data except the last
12 months for testing
    train_data, test_data = train_test_split(share_price,
train_size=train_size)

    # Perform hyperparameter tuning for Auto ARIMA
    model = pm.auto_arima(
        train_data,
        seasonal=True,
        m=12,
        stepwise=True,

```

```

        error_action='ignore', # Ignore orders that don't converge
        suppress_warnings=True, # Suppress warnings
        scoring='mse', # Use mean squared error for model selection
        max_order=None, # Maximum ARIMA order
        trace=True) # Print diagnostic information)
# Make forecasts for the next 12 months
forecast, conf_int = model.predict(n_periods=12,
return_conf_int=True)

# Extract the actual values for the test period
actual_values = test_data['Close'].values

# Calculate MAPE for the 12-month period
mape = calculate_mape(actual_values, forecast)

# Append results to the DataFrame
mape_results.append({'Company': company_name, 'Sector': sector,
'MAPE': mape})

# Print the model results
print(f"Company: {company_name}")
print(f"Sector: {sector}")
print(f"Best Model Order: {model.order}")
print(f"Forecast: {forecast}")
print(f"Confidence Interval: {conf_int}")
print(f"Actual Values: {actual_values}")
print(f"MAPE: {mape}")
print("\n")

# Print the results for the top 2 companies in each sector
for result in mape_results:
    print(result)

```

Performing stepwise search to minimize aic

```

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=8093.653, Time=4.38 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=8103.411, Time=0.09 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=8094.557, Time=0.73 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=8095.065, Time=0.65 sec
ARIMA(0,1,0)(0,0,0)[12] : AIC=8103.188, Time=0.06 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=8092.344, Time=3.99 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=8091.324, Time=1.49 sec
ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=8092.273, Time=4.05 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=8094.921, Time=0.75 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=8094.922, Time=0.66 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=inf, Time=2.83 sec
ARIMA(2,1,3)(0,0,0)[12] intercept : AIC=8088.346, Time=1.49 sec
ARIMA(2,1,3)(1,0,0)[12] intercept : AIC=8089.320, Time=3.81 sec
ARIMA(2,1,3)(0,0,1)[12] intercept : AIC=8089.390, Time=4.00 sec
ARIMA(2,1,3)(1,0,1)[12] intercept : AIC=8090.664, Time=8.49 sec
ARIMA(1,1,3)(0,0,0)[12] intercept : AIC=8086.551, Time=0.87 sec

```

```

ARIMA(1,1,3)(1,0,0)[12] intercept : AIC=8087.520, Time=1.95 sec
ARIMA(1,1,3)(0,0,1)[12] intercept : AIC=8087.590, Time=2.60 sec
ARIMA(1,1,3)(1,0,1)[12] intercept : AIC=8088.855, Time=3.00 sec
ARIMA(0,1,3)(0,0,0)[12] intercept : AIC=8087.727, Time=0.28 sec
ARIMA(1,1,4)(0,0,0)[12] intercept : AIC=8089.236, Time=1.15 sec
ARIMA(0,1,2)(0,0,0)[12] intercept : AIC=8095.674, Time=0.23 sec
ARIMA(0,1,4)(0,0,0)[12] intercept : AIC=8088.848, Time=0.38 sec
ARIMA(2,1,4)(0,0,0)[12] intercept : AIC=8090.269, Time=1.43 sec
ARIMA(1,1,3)(0,0,0)[12] : AIC=8087.176, Time=0.38 sec

```

Best model: ARIMA(1,1,3)(0,0,0)[12] intercept

Total fit time: 49.838 seconds

Company: Alphabet Inc. (Class C)

Sector: Communication Services

Best Model Order: (1, 1, 3)

Forecast: 2505 98.433335

2506 98.566411

2507 98.646336

2508 98.703488

2509 98.748827

2510 98.788038

2511 98.824069

2512 98.858450

2513 98.891976

2514 98.925058

2515 98.957909

2516 98.990642

dtype: float64

Confidence Interval: [[96.05521007 100.8114599]

[95.31588945 101.81693237]

[94.69026716 102.60240451]

[94.21738904 103.18958795]

[93.81946607 103.67818883]

[93.46637919 104.10969628]

[93.14395262 104.50418461]

[92.84452164 104.87237856]

[92.56344869 105.22050313]

[92.2976554 105.55246025]

[92.04494222 105.87087669]

[91.80364711 106.17763611]]

Actual Values: [99.70999908 97.18000031 100.77999878 101.38999939
100.29000092

100.52999878 101.48000336 102.97000122 104.93000031 94.81999969

92.59999847 96.58000183]

MAPE: 0.030665961117357168

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=8113.243, Time=3.77 sec

ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=8123.896, Time=0.03 sec

ARIMA(1,1,0)(1,0,0)[12] intercept	: AIC=8113.854, Time=0.46 sec
ARIMA(0,1,1)(0,0,1)[12] intercept	: AIC=8114.333, Time=0.65 sec
ARIMA(0,1,0)(0,0,0)[12]	: AIC=8123.624, Time=0.07 sec
ARIMA(2,1,2)(0,0,1)[12] intercept	: AIC=8111.974, Time=3.97 sec
ARIMA(2,1,2)(0,0,0)[12] intercept	: AIC=8110.778, Time=1.31 sec
ARIMA(2,1,2)(1,0,0)[12] intercept	: AIC=8111.914, Time=3.56 sec
ARIMA(1,1,2)(0,0,0)[12] intercept	: AIC=8114.730, Time=0.78 sec
ARIMA(2,1,1)(0,0,0)[12] intercept	: AIC=8114.733, Time=1.23 sec
ARIMA(3,1,2)(0,0,0)[12] intercept	: AIC=inf, Time=3.03 sec
ARIMA(2,1,3)(0,0,0)[12] intercept	: AIC=8108.887, Time=1.32 sec
ARIMA(2,1,3)(1,0,0)[12] intercept	: AIC=8110.067, Time=3.64 sec
ARIMA(2,1,3)(0,0,1)[12] intercept	: AIC=8110.124, Time=4.00 sec
ARIMA(2,1,3)(1,0,1)[12] intercept	: AIC=8111.363, Time=12.42 sec
ARIMA(1,1,3)(0,0,0)[12] intercept	: AIC=8107.006, Time=1.12 sec
ARIMA(1,1,3)(1,0,0)[12] intercept	: AIC=8108.188, Time=2.56 sec
ARIMA(1,1,3)(0,0,1)[12] intercept	: AIC=8108.245, Time=3.12 sec
ARIMA(1,1,3)(1,0,1)[12] intercept	: AIC=8109.470, Time=4.32 sec
ARIMA(0,1,3)(0,0,0)[12] intercept	: AIC=8108.508, Time=0.38 sec
ARIMA(1,1,4)(0,0,0)[12] intercept	: AIC=8109.632, Time=1.28 sec
ARIMA(0,1,2)(0,0,0)[12] intercept	: AIC=8114.943, Time=0.37 sec
ARIMA(0,1,4)(0,0,0)[12] intercept	: AIC=8109.674, Time=0.56 sec
ARIMA(2,1,4)(0,0,0)[12] intercept	: AIC=8110.952, Time=1.77 sec
ARIMA(1,1,3)(0,0,0)[12]	: AIC=8107.590, Time=0.91 sec

Best model: ARIMA(1,1,3)(0,0,0)[12] intercept

Total fit time: 56.635 seconds

Company: Alphabet Inc. (Class A)

Sector: Communication Services

Best Model Order: (1, 1, 3)

Forecast: 2505 97.679485

2506 97.813768

2507 97.889974

2508 97.946538

2509 97.992242

2510 98.031943

2511 98.068325

2512 98.102871

2513 98.136404

2514 98.169375

2515 98.202036

2516 98.234526

dtype: float64

Confidence Interval: [[95.29162128 100.06734817]

[94.55691551 101.07061996]

[93.93211597 101.84783189]

[93.4568124 102.43626321]

[93.05738195 102.92710267]

[92.70392184 103.35996444]

[92.38189178 103.75475782]


```

[ 92.08328674 104.12245613]
[ 91.80323995 104.46956719]
[ 91.53853693 104.80021271]
[ 91.28690407 105.11716792]
[ 91.0466427 105.42240877]]
Actual Values: [ 99.05999756 96.55999756 99.97000122 100.76999664
99.62999725
99.97000122 101.12999725 102.51999664 104.48000336 94.93000031
92.22000122 96.29000092]
MAPE: 0.030402662356181276

```

Performing stepwise search to minimize aic

```

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=inf, Time=8.66 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=10359.292, Time=0.05 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=10360.298, Time=0.57 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=10360.308, Time=0.64 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=10358.410, Time=0.05 sec
ARIMA(0,1,0)(1,0,0)[12] intercept : AIC=10360.083, Time=0.37 sec
ARIMA(0,1,0)(0,0,1)[12] intercept : AIC=10360.016, Time=0.34 sec
ARIMA(0,1,0)(1,0,1)[12] intercept : AIC=10361.625, Time=1.12 sec
ARIMA(1,1,0)(0,0,0)[12] intercept : AIC=10359.371, Time=0.15 sec
ARIMA(0,1,1)(0,0,0)[12] intercept : AIC=10359.453, Time=0.18 sec
ARIMA(1,1,1)(0,0,0)[12] intercept : AIC=10356.676, Time=0.62 sec
ARIMA(1,1,1)(1,0,0)[12] intercept : AIC=10357.692, Time=2.04 sec
ARIMA(1,1,1)(0,0,1)[12] intercept : AIC=10357.638, Time=1.85 sec
ARIMA(1,1,1)(1,0,1)[12] intercept : AIC=10359.285, Time=3.30 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=10357.443, Time=0.78 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=10357.259, Time=0.58 sec
ARIMA(0,1,2)(0,0,0)[12] intercept : AIC=10360.373, Time=0.24 sec
ARIMA(2,1,0)(0,0,0)[12] intercept : AIC=10360.012, Time=0.17 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=10335.394, Time=2.30 sec
ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=10336.720, Time=6.61 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=10336.672, Time=6.14 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=10342.010, Time=2.90 sec
ARIMA(2,1,3)(0,0,0)[12] intercept : AIC=10336.692, Time=2.86 sec
ARIMA(1,1,3)(0,0,0)[12] intercept : AIC=10354.640, Time=0.89 sec
ARIMA(3,1,1)(0,0,0)[12] intercept : AIC=10355.258, Time=0.85 sec
ARIMA(3,1,3)(0,0,0)[12] intercept : AIC=inf, Time=3.60 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=10334.521, Time=1.25 sec
ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=10335.875, Time=3.06 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=10335.840, Time=2.75 sec
ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=inf, Time=4.03 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=10356.373, Time=0.30 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=10356.563, Time=0.33 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=10335.853, Time=1.40 sec
ARIMA(2,1,3)(0,0,0)[12] intercept : AIC=10335.858, Time=1.52 sec
ARIMA(1,1,1)(0,0,0)[12] intercept : AIC=10355.849, Time=0.28 sec
ARIMA(1,1,3)(0,0,0)[12] intercept : AIC=10353.867, Time=0.34 sec

```

```
ARIMA(3,1,1)(0,0,0)[12] : AIC=10354.485, Time=0.37 sec
ARIMA(3,1,3)(0,0,0)[12] : AIC=inf, Time=1.87 sec
```

Best model: ARIMA(2,1,2)(0,0,0)[12]

Total fit time: 65.401 seconds

Company: Amazon

Sector: Consumer Discretionary

Best Model Order: (2, 1, 2)

Forecast: 2505 113.179334

2506 112.922223

2507 113.089677

2508 113.053107

2509 112.953777

2510 113.157779

2511 112.906992

2512 113.136335

2513 112.988530

2514 113.018493

2515 113.110102

2516 112.925587

dtype: float64

Confidence Interval: [[109.45275717 116.90591037]

[107.65343555 118.19101065]

[106.60372921 119.57562437]

[105.59291516 120.51329869]

[104.58421056 121.32334254]

[104.00540647 122.31015154]

[103.01495585 122.7990274]

[102.5551654 123.71750397]

[101.7780345 124.19902611]

[101.18458427 124.85240142]

[100.71320659 125.506998]

[99.9669495 125.88422409]]

Actual Values: [112.52999878 106.90000153 113.79000092 116.36000061

115.06999969

115.25 119.31999969 119.81999969 120.59999847 115.66000366

110.95999908 103.41000366]

MAPE: 0.036674863194271

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=14925.577, Time=8.14 sec

ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=14970.347, Time=0.07 sec

ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=14967.013, Time=0.82 sec

ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=14967.957, Time=0.75 sec

ARIMA(0,1,0)(0,0,0)[12] : AIC=14969.150, Time=0.05 sec

ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=14973.104, Time=2.37 sec

ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=14972.132, Time=2.77 sec

ARIMA(2,1,2)(2,0,1)[12] intercept : AIC=14932.858, Time=12.56 sec

ARIMA(2,1,2)(1,0,2)[12] intercept : AIC=14912.687, Time=23.90 sec

ARIMA(2,1,2)(0,0,2)[12]	intercept	: AIC=14930.032, Time=8.44 sec
ARIMA(2,1,2)(2,0,2)[12]	intercept	: AIC=14930.601, Time=20.07 sec
ARIMA(1,1,2)(1,0,2)[12]	intercept	: AIC=14928.837, Time=11.02 sec
ARIMA(2,1,1)(1,0,2)[12]	intercept	: AIC=14926.990, Time=14.82 sec
ARIMA(3,1,2)(1,0,2)[12]	intercept	: AIC=14887.123, Time=20.69 sec
ARIMA(3,1,2)(0,0,2)[12]	intercept	: AIC=14886.227, Time=19.99 sec
ARIMA(3,1,2)(0,0,1)[12]	intercept	: AIC=14920.464, Time=6.92 sec
ARIMA(3,1,2)(1,0,1)[12]	intercept	: AIC=14913.464, Time=7.59 sec
ARIMA(3,1,1)(0,0,2)[12]	intercept	: AIC=14928.222, Time=10.31 sec
ARIMA(4,1,2)(0,0,2)[12]	intercept	: AIC=14881.648, Time=20.38 sec
ARIMA(4,1,2)(0,0,1)[12]	intercept	: AIC=14913.846, Time=8.04 sec
ARIMA(4,1,2)(1,0,2)[12]	intercept	: AIC=14883.077, Time=22.75 sec
ARIMA(4,1,2)(1,0,1)[12]	intercept	: AIC=14907.897, Time=10.19 sec
ARIMA(4,1,1)(0,0,2)[12]	intercept	: AIC=14925.250, Time=9.09 sec
ARIMA(5,1,2)(0,0,2)[12]	intercept	: AIC=14878.780, Time=23.05 sec
ARIMA(5,1,2)(0,0,1)[12]	intercept	: AIC=14915.808, Time=8.49 sec
ARIMA(5,1,2)(1,0,2)[12]	intercept	: AIC=14878.328, Time=25.26 sec
ARIMA(5,1,2)(1,0,1)[12]	intercept	: AIC=14911.393, Time=13.01 sec
ARIMA(5,1,2)(2,0,2)[12]	intercept	: AIC=14887.654, Time=27.46 sec
ARIMA(5,1,2)(2,0,1)[12]	intercept	: AIC=14880.780, Time=29.65 sec
ARIMA(5,1,1)(1,0,2)[12]	intercept	: AIC=14926.605, Time=15.88 sec
ARIMA(5,1,3)(1,0,2)[12]	intercept	: AIC=14886.217, Time=31.71 sec
ARIMA(4,1,1)(1,0,2)[12]	intercept	: AIC=14925.798, Time=12.01 sec
ARIMA(4,1,3)(1,0,2)[12]	intercept	: AIC=14884.034, Time=27.51 sec
ARIMA(5,1,2)(1,0,2)[12]		: AIC=14876.669, Time=10.16 sec
ARIMA(5,1,2)(0,0,2)[12]		: AIC=14877.238, Time=8.21 sec
ARIMA(5,1,2)(1,0,1)[12]		: AIC=14908.123, Time=5.65 sec
ARIMA(5,1,2)(2,0,2)[12]		: AIC=14885.704, Time=13.52 sec
ARIMA(5,1,2)(0,0,1)[12]		: AIC=14929.780, Time=4.59 sec
ARIMA(5,1,2)(2,0,1)[12]		: AIC=14879.098, Time=10.93 sec
ARIMA(4,1,2)(1,0,2)[12]		: AIC=14881.623, Time=9.52 sec
ARIMA(5,1,1)(1,0,2)[12]		: AIC=14925.066, Time=5.00 sec
ARIMA(5,1,3)(1,0,2)[12]		: AIC=14884.531, Time=10.90 sec
ARIMA(4,1,1)(1,0,2)[12]		: AIC=14924.235, Time=4.56 sec
ARIMA(4,1,3)(1,0,2)[12]		: AIC=14882.578, Time=12.51 sec

Best model: ARIMA(5,1,2)(1,0,2)[12]

Total fit time: 551.356 seconds

Company: Tesla Inc.

Sector: Consumer Discretionary

Best Model Order: (5, 1, 2)

Forecast: 2505 221.722821

2506 219.872360

2507 219.055716

2508 218.281069

2509 219.139586

2510 217.692822

2511 218.497478

2512 218.270364

```

2513    215.806090
2514    214.716942
2515    211.698188
2516    212.946313
dtype: float64
Confidence Interval: [[212.51399089 230.93165113]
 [206.94557479 232.79914477]
 [203.20445005 234.90698195]
 [199.898926    236.66321132]
 [198.23344     240.04573257]
 [194.54919362 240.83644944]
 [193.46113227 243.53382429]
 [191.24830755 245.29242018]
 [187.13689377 244.4752864 ]
 [184.31700037 245.11688421]
 [179.79547085 243.60090476]
 [179.51883666 246.37378943]]
Actual Values: [221.72000122 204.99000549 219.3500061  220.19000244
222.03999329
 207.27999878 214.44000244 211.25          222.41999817 224.63999939
 225.08999634 228.52000427]
MAPE: 0.03330257792758643

```

Performing stepwise search to minimize aic

```

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=inf, Time=6.75 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=8835.918, Time=0.05 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=8757.384, Time=0.48 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=8764.971, Time=0.46 sec
ARIMA(0,1,0)(0,0,0)[12]          : AIC=8835.928, Time=0.05 sec
ARIMA(1,1,0)(0,0,0)[12] intercept : AIC=8756.488, Time=0.13 sec
ARIMA(1,1,0)(0,0,1)[12] intercept : AIC=8757.401, Time=0.42 sec
ARIMA(1,1,0)(1,0,1)[12] intercept : AIC=8759.328, Time=2.33 sec
ARIMA(2,1,0)(0,0,0)[12] intercept : AIC=8753.037, Time=0.21 sec
ARIMA(2,1,0)(1,0,0)[12] intercept : AIC=8754.182, Time=0.77 sec
ARIMA(2,1,0)(0,0,1)[12] intercept : AIC=8754.202, Time=0.58 sec
ARIMA(2,1,0)(1,0,1)[12] intercept : AIC=8756.062, Time=2.92 sec
ARIMA(3,1,0)(0,0,0)[12] intercept : AIC=8748.689, Time=0.23 sec
ARIMA(3,1,0)(1,0,0)[12] intercept : AIC=8750.183, Time=0.73 sec
ARIMA(3,1,0)(0,0,1)[12] intercept : AIC=8750.197, Time=0.59 sec
ARIMA(3,1,0)(1,0,1)[12] intercept : AIC=8752.030, Time=2.96 sec
ARIMA(4,1,0)(0,0,0)[12] intercept : AIC=8738.124, Time=0.28 sec
ARIMA(4,1,0)(1,0,0)[12] intercept : AIC=8739.590, Time=0.97 sec
ARIMA(4,1,0)(0,0,1)[12] intercept : AIC=8739.599, Time=0.70 sec
ARIMA(4,1,0)(1,0,1)[12] intercept : AIC=8741.525, Time=4.88 sec
ARIMA(5,1,0)(0,0,0)[12] intercept : AIC=8736.219, Time=0.38 sec
ARIMA(5,1,0)(1,0,0)[12] intercept : AIC=8737.416, Time=1.09 sec
ARIMA(5,1,0)(0,0,1)[12] intercept : AIC=8737.432, Time=0.82 sec
ARIMA(5,1,0)(1,0,1)[12] intercept : AIC=8739.337, Time=3.51 sec

```

ARIMA(5,1,1)(0,0,0)[12] intercept	: AIC=8701.505, Time=1.73 sec
ARIMA(5,1,1)(1,0,0)[12] intercept	: AIC=8703.355, Time=4.46 sec
ARIMA(5,1,1)(0,0,1)[12] intercept	: AIC=8703.359, Time=4.00 sec
ARIMA(5,1,1)(1,0,1)[12] intercept	: AIC=inf, Time=11.49 sec
ARIMA(4,1,1)(0,0,0)[12] intercept	: AIC=8699.838, Time=1.28 sec
ARIMA(4,1,1)(1,0,0)[12] intercept	: AIC=8701.751, Time=3.75 sec
ARIMA(4,1,1)(0,0,1)[12] intercept	: AIC=8701.753, Time=3.89 sec
ARIMA(4,1,1)(1,0,1)[12] intercept	: AIC=8703.791, Time=4.60 sec
ARIMA(3,1,1)(0,0,0)[12] intercept	: AIC=8745.394, Time=1.28 sec
ARIMA(4,1,2)(0,0,0)[12] intercept	: AIC=8696.833, Time=2.91 sec
ARIMA(4,1,2)(1,0,0)[12] intercept	: AIC=8695.593, Time=8.36 sec
ARIMA(4,1,2)(2,0,0)[12] intercept	: AIC=8697.061, Time=19.27 sec
ARIMA(4,1,2)(1,0,1)[12] intercept	: AIC=8697.359, Time=9.38 sec
ARIMA(4,1,2)(0,0,1)[12] intercept	: AIC=8695.741, Time=7.44 sec
ARIMA(4,1,2)(2,0,1)[12] intercept	: AIC=8698.709, Time=24.67 sec
ARIMA(3,1,2)(1,0,0)[12] intercept	: AIC=8739.621, Time=4.34 sec
ARIMA(5,1,2)(1,0,0)[12] intercept	: AIC=inf, Time=9.60 sec
ARIMA(4,1,3)(1,0,0)[12] intercept	: AIC=8694.937, Time=10.14 sec
ARIMA(4,1,3)(0,0,0)[12] intercept	: AIC=8693.580, Time=3.67 sec
ARIMA(4,1,3)(0,0,1)[12] intercept	: AIC=8696.549, Time=9.80 sec
ARIMA(4,1,3)(1,0,1)[12] intercept	: AIC=8701.182, Time=10.22 sec
ARIMA(3,1,3)(0,0,0)[12] intercept	: AIC=8697.217, Time=2.30 sec
ARIMA(5,1,3)(0,0,0)[12] intercept	: AIC=8702.333, Time=4.06 sec
ARIMA(4,1,4)(0,0,0)[12] intercept	: AIC=8683.894, Time=2.64 sec
ARIMA(4,1,4)(1,0,0)[12] intercept	: AIC=8685.861, Time=8.97 sec
ARIMA(4,1,4)(0,0,1)[12] intercept	: AIC=8685.862, Time=10.75 sec
ARIMA(4,1,4)(1,0,1)[12] intercept	: AIC=8687.894, Time=10.46 sec
ARIMA(3,1,4)(0,0,0)[12] intercept	: AIC=inf, Time=3.52 sec
ARIMA(5,1,4)(0,0,0)[12] intercept	: AIC=8685.878, Time=4.06 sec
ARIMA(4,1,5)(0,0,0)[12] intercept	: AIC=8685.879, Time=4.94 sec
ARIMA(3,1,5)(0,0,0)[12] intercept	: AIC=8688.055, Time=4.72 sec
ARIMA(5,1,5)(0,0,0)[12] intercept	: AIC=8687.887, Time=5.20 sec
ARIMA(4,1,4)(0,0,0)[12]	: AIC=8684.813, Time=2.25 sec

Best model: ARIMA(4,1,4)(0,0,0)[12] intercept

Total fit time: 252.544 seconds

Company: PepsiCo

Sector: Consumer Staples

Best Model Order: (4, 1, 4)

Forecast: 2505 168.534217

2506 169.046263

2507 168.326273

2508 168.041304

2509 168.155699

2510 167.865412

2511 168.668054

2512 168.493811

2513 169.132302

2514 168.872886

```

2515    168.913552
2516    168.727024
dtype: float64
Confidence Interval: [[165.8592364  171.20919668]
 [165.54234491 172.55018118]
 [164.07238212 172.58016296]
 [163.25245328 172.83015528]
 [162.92744078 173.38395697]
 [162.18185444 173.54896984]
 [162.61279679 174.72331127]
 [161.98406803 175.00355448]
 [162.24474084 176.0198626 ]
 [161.56816662 176.17760597]
 [161.25517895 176.57192457]
 [160.73434657 176.71970144]]
Actual Values: [174.61000061 170.19000244 172.72999573 175.05999756
173.36000061
171.46000671 173.05999756 177.67999268 178.27000427 179.07000732
178.88000488 182.22999573]
MAPE: 0.039432099198532126

```

```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=14221.908, Time=7.00 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=14223.972, Time=0.03 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=14225.995, Time=0.54 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=14225.926, Time=0.76 sec
ARIMA(0,1,0)(0,0,0)[12] : AIC=14225.126, Time=0.03 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=14213.926, Time=6.17 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=14212.310, Time=2.28 sec
ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=14213.966, Time=6.08 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=14226.990, Time=0.76 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=14226.974, Time=0.74 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=14216.502, Time=2.33 sec
ARIMA(2,1,3)(0,0,0)[12] intercept : AIC=14192.275, Time=3.19 sec
ARIMA(2,1,3)(1,0,0)[12] intercept : AIC=14194.223, Time=8.04 sec
ARIMA(2,1,3)(0,0,1)[12] intercept : AIC=14194.184, Time=8.22 sec
ARIMA(2,1,3)(1,0,1)[12] intercept : AIC=14194.473, Time=10.76 sec
ARIMA(1,1,3)(0,0,0)[12] intercept : AIC=14209.746, Time=1.15 sec
ARIMA(3,1,3)(0,0,0)[12] intercept : AIC=14218.577, Time=3.27 sec
ARIMA(2,1,4)(0,0,0)[12] intercept : AIC=14218.487, Time=4.16 sec
ARIMA(1,1,4)(0,0,0)[12] intercept : AIC=14202.708, Time=1.43 sec
ARIMA(3,1,4)(0,0,0)[12] intercept : AIC=14220.522, Time=2.81 sec
ARIMA(2,1,3)(0,0,0)[12] : AIC=14195.794, Time=1.52 sec

```

```

Best model: ARIMA(2,1,3)(0,0,0)[12] intercept
Total fit time: 71.277 seconds
Company: Costco
Sector: Consumer Staples
Best Model Order: (2, 1, 3)

```

```

Forecast: 2505      467.433322
2506      466.644992
2507      467.420528
2508      467.287536
2509      467.342309
2510      467.896198
2511      467.384818
2512      468.320065
2513      467.646895
2514      468.521045
2515      468.105840
2516      468.574747
dtype: float64
Confidence Interval: [[459.37617253 475.49047203]
 [455.30320719 477.98677718]
 [453.38851696 481.45253902]
 [451.06643     483.50864229]
 [449.21000528 485.47461242]
 [447.96295979 487.82943538]
 [445.90773978 488.86189563]
 [445.28010777 491.36002309]
 [443.26271413 492.03107529]
 [442.76133395 494.28075609]
 [441.11483457 495.09684538]
 [440.36443305 496.78505996]]
Actual Values: [467.98999023 454.6499939 464.17001343 473.26998901
471.42999268
464.61999512 478.17999268 496.97000122 499.05999756 499.45001221
496.54000854 510.86999512]
MAPE: 0.03400739383792633

```

```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=5123.957, Time=7.62 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=5142.095, Time=0.16 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=5134.369, Time=0.85 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=5134.657, Time=0.89 sec
ARIMA(0,1,0)(0,0,0)[12]          : AIC=5140.985, Time=0.08 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=5121.980, Time=4.37 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=5120.547, Time=1.76 sec
ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=5121.987, Time=4.58 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=5127.924, Time=0.92 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=5127.918, Time=2.59 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=5118.454, Time=3.09 sec
ARIMA(3,1,2)(1,0,0)[12] intercept : AIC=5134.014, Time=9.08 sec
ARIMA(3,1,2)(0,0,1)[12] intercept : AIC=5134.005, Time=9.25 sec
ARIMA(3,1,2)(1,0,1)[12] intercept : AIC=5136.005, Time=8.25 sec
ARIMA(3,1,1)(0,0,0)[12] intercept : AIC=5123.558, Time=1.27 sec
ARIMA(4,1,2)(0,0,0)[12] intercept : AIC=5117.443, Time=3.30 sec

```

```

ARIMA(4,1,2)(1,0,0)[12] intercept : AIC=5118.257, Time=9.66 sec
ARIMA(4,1,2)(0,0,1)[12] intercept : AIC=5117.420, Time=7.53 sec
ARIMA(4,1,2)(1,0,1)[12] intercept : AIC=5119.769, Time=9.58 sec
ARIMA(4,1,2)(0,0,2)[12] intercept : AIC=5120.906, Time=21.58 sec
ARIMA(4,1,2)(1,0,2)[12] intercept : AIC=5123.754, Time=20.77 sec
ARIMA(4,1,1)(0,0,1)[12] intercept : AIC=5123.610, Time=3.37 sec
ARIMA(5,1,2)(0,0,1)[12] intercept : AIC=5117.844, Time=8.42 sec
ARIMA(4,1,3)(0,0,1)[12] intercept : AIC=5124.135, Time=9.95 sec
ARIMA(3,1,1)(0,0,1)[12] intercept : AIC=5124.665, Time=3.73 sec
ARIMA(3,1,3)(0,0,1)[12] intercept : AIC=5117.260, Time=8.50 sec
ARIMA(3,1,3)(0,0,0)[12] intercept : AIC=5117.194, Time=2.87 sec
ARIMA(3,1,3)(1,0,0)[12] intercept : AIC=5117.520, Time=6.76 sec
ARIMA(3,1,3)(1,0,1)[12] intercept : AIC=5119.492, Time=9.36 sec
ARIMA(2,1,3)(0,0,0)[12] intercept : AIC=5122.464, Time=1.92 sec
ARIMA(4,1,3)(0,0,0)[12] intercept : AIC=5122.893, Time=3.58 sec
ARIMA(3,1,4)(0,0,0)[12] intercept : AIC=inf, Time=3.95 sec
ARIMA(2,1,4)(0,0,0)[12] intercept : AIC=5125.017, Time=1.44 sec
ARIMA(4,1,4)(0,0,0)[12] intercept : AIC=5120.708, Time=4.05 sec
ARIMA(3,1,3)(0,0,0)[12] intercept : AIC=5117.375, Time=1.45 sec

```

Best model: ARIMA(3,1,3)(0,0,0)[12] intercept

Total fit time: 196.546 seconds

Company: AstraZeneca

Sector: Health Care

Best Model Order: (3, 1, 3)

Forecast: 2505 55.182048

2506 55.409210

2507 55.406668

2508 55.682784

2509 55.655097

2510 55.904411

2511 55.911417

2512 56.094968

2513 56.157715

2514 56.272719

2515 56.381646

2516 56.448292

dtype: float64

Confidence Interval: [[53.86840361 56.49569232]

[53.60441012 57.21401042]

[53.22156421 57.59177205]

[53.17799886 58.1875698]

[52.89128798 58.41890652]

[52.89218503 58.91663612]

[52.69615716 59.1266774]

[52.67702452 59.51291172]

[52.56847238 59.74695823]

[52.51410932 60.03132813]

[52.47251459 60.29077739]

[52.39465847 60.50192543]]

Actual Values: [55.52999878 54.97000122 56.18000031 55.97000122
54.50999832 54.34999847
54.97000122 55.18000031 55.90000153 57.95999908 57.61000061
58.70999908]
MAPE: 0.017502849181341738

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12] intercept	: AIC=12394.998, Time=6.22 sec
ARIMA(0,1,0)(0,0,0)[12] intercept	: AIC=12412.633, Time=0.05 sec
ARIMA(1,1,0)(1,0,0)[12] intercept	: AIC=12401.058, Time=0.52 sec
ARIMA(0,1,1)(0,0,1)[12] intercept	: AIC=12401.294, Time=0.64 sec
ARIMA(0,1,0)(0,0,0)[12]	: AIC=12411.862, Time=0.05 sec
ARIMA(2,1,2)(0,0,1)[12] intercept	: AIC=12395.142, Time=6.66 sec
ARIMA(2,1,2)(1,0,0)[12] intercept	: AIC=12392.489, Time=5.88 sec
ARIMA(2,1,2)(0,0,0)[12] intercept	: AIC=inf, Time=2.25 sec
ARIMA(2,1,2)(2,0,0)[12] intercept	: AIC=12396.417, Time=15.47 sec
ARIMA(2,1,2)(2,0,1)[12] intercept	: AIC=12400.722, Time=18.25 sec
ARIMA(1,1,2)(1,0,0)[12] intercept	: AIC=12404.621, Time=1.80 sec
ARIMA(2,1,1)(1,0,0)[12] intercept	: AIC=12404.648, Time=2.52 sec
ARIMA(3,1,2)(1,0,0)[12] intercept	: AIC=12353.477, Time=7.42 sec
ARIMA(3,1,2)(0,0,0)[12] intercept	: AIC=12351.532, Time=2.69 sec
ARIMA(3,1,2)(0,0,1)[12] intercept	: AIC=12353.452, Time=7.09 sec
ARIMA(3,1,2)(1,0,1)[12] intercept	: AIC=12355.354, Time=8.67 sec
ARIMA(3,1,1)(0,0,0)[12] intercept	: AIC=12403.829, Time=0.77 sec
ARIMA(4,1,2)(0,0,0)[12] intercept	: AIC=12346.677, Time=3.62 sec
ARIMA(4,1,2)(1,0,0)[12] intercept	: AIC=12353.416, Time=9.58 sec
ARIMA(4,1,2)(0,0,1)[12] intercept	: AIC=12348.448, Time=7.88 sec
ARIMA(4,1,2)(1,0,1)[12] intercept	: AIC=12350.638, Time=11.35 sec
ARIMA(4,1,1)(0,0,0)[12] intercept	: AIC=12384.703, Time=1.44 sec
ARIMA(5,1,2)(0,0,0)[12] intercept	: AIC=12348.163, Time=4.14 sec
ARIMA(4,1,3)(0,0,0)[12] intercept	: AIC=12342.502, Time=4.06 sec
ARIMA(4,1,3)(1,0,0)[12] intercept	: AIC=12344.530, Time=10.66 sec
ARIMA(4,1,3)(0,0,1)[12] intercept	: AIC=12344.229, Time=11.94 sec
ARIMA(4,1,3)(1,0,1)[12] intercept	: AIC=12346.261, Time=11.30 sec
ARIMA(3,1,3)(0,0,0)[12] intercept	: AIC=12344.070, Time=3.23 sec
ARIMA(5,1,3)(0,0,0)[12] intercept	: AIC=12339.118, Time=4.31 sec
ARIMA(5,1,3)(1,0,0)[12] intercept	: AIC=12342.297, Time=10.81 sec
ARIMA(5,1,3)(0,0,1)[12] intercept	: AIC=inf, Time=11.52 sec
ARIMA(5,1,3)(1,0,1)[12] intercept	: AIC=12342.295, Time=12.93 sec
ARIMA(5,1,4)(0,0,0)[12] intercept	: AIC=inf, Time=4.75 sec
ARIMA(4,1,4)(0,0,0)[12] intercept	: AIC=12342.341, Time=4.73 sec
ARIMA(5,1,3)(0,0,0)[12]	: AIC=12342.689, Time=3.00 sec

Best model: ARIMA(5,1,3)(0,0,0)[12] intercept
Total fit time: 218.206 seconds
Company: Amgen
Sector: Health Care
Best Model Order: (5, 1, 3)
Forecast: 2505 245.997694

```

2506      245.879157
2507      245.935349
2508      246.057832
2509      245.948816
2510      246.150557
2511      246.030474
2512      246.180712
2513      246.168144
2514      246.185582
2515      246.307624
2516      246.221692
dtype: float64
Confidence Interval: [[240.45986003 251.53552866]
 [238.24071568 253.51759881]
 [236.70211802 255.16857938]
 [235.59217682 256.52348655]
 [234.38403901 257.51359222]
 [233.58072017 258.72039293]
 [232.61840346 259.44254455]
 [231.91798475 260.4434401 ]
 [231.2010974  261.13519103]
 [230.50340923 261.86775477]
 [229.99362657 262.62162118]
 [229.31000736 263.13337626]]
Actual Values: [251.66000366 251.33999634 252.92999268 252.11999512
248.19000244
 247.44999695 251.94000244 261.32000732 259.98999023 266.66000366
 267.23001099 273.80999756]
MAPE: 0.04171600800666341

```

```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=10753.745, Time=5.16 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=10753.281, Time=0.03 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=10749.333, Time=0.56 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=10749.405, Time=0.67 sec
ARIMA(0,1,0)(0,0,0)[12] : AIC=10752.500, Time=0.03 sec
ARIMA(1,1,0)(0,0,0)[12] intercept : AIC=10750.211, Time=0.14 sec
ARIMA(1,1,0)(2,0,0)[12] intercept : AIC=10751.329, Time=1.48 sec
ARIMA(1,1,0)(1,0,1)[12] intercept : AIC=10747.934, Time=2.05 sec
ARIMA(1,1,0)(0,0,1)[12] intercept : AIC=10749.317, Time=0.48 sec
ARIMA(1,1,0)(2,0,1)[12] intercept : AIC=10749.717, Time=7.11 sec
ARIMA(1,1,0)(1,0,2)[12] intercept : AIC=10753.316, Time=2.38 sec
ARIMA(1,1,0)(0,0,2)[12] intercept : AIC=10751.271, Time=1.60 sec
ARIMA(1,1,0)(2,0,2)[12] intercept : AIC=inf, Time=8.99 sec
ARIMA(0,1,0)(1,0,1)[12] intercept : AIC=10750.654, Time=1.67 sec
ARIMA(2,1,0)(1,0,1)[12] intercept : AIC=10749.748, Time=4.27 sec
ARIMA(1,1,1)(1,0,1)[12] intercept : AIC=10749.749, Time=6.50 sec
ARIMA(0,1,1)(1,0,1)[12] intercept : AIC=10748.020, Time=2.34 sec

```

```

ARIMA(2,1,1)(1,0,1)[12] intercept : AIC=10717.883, Time=8.51 sec
ARIMA(2,1,1)(0,0,1)[12] intercept : AIC=10717.730, Time=4.10 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=10720.903, Time=1.26 sec
ARIMA(2,1,1)(0,0,2)[12] intercept : AIC=10719.694, Time=10.73 sec
ARIMA(2,1,1)(1,0,0)[12] intercept : AIC=10717.751, Time=4.11 sec
ARIMA(2,1,1)(1,0,2)[12] intercept : AIC=10720.596, Time=18.48 sec
ARIMA(1,1,1)(0,0,1)[12] intercept : AIC=10751.153, Time=1.31 sec
ARIMA(2,1,0)(0,0,1)[12] intercept : AIC=10751.138, Time=0.71 sec
ARIMA(3,1,1)(0,0,1)[12] intercept : AIC=10755.115, Time=1.09 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=10755.140, Time=2.09 sec
ARIMA(1,1,2)(0,0,1)[12] intercept : AIC=10753.144, Time=1.07 sec
ARIMA(3,1,0)(0,0,1)[12] intercept : AIC=10753.119, Time=0.74 sec
ARIMA(3,1,2)(0,0,1)[12] intercept : AIC=10757.058, Time=2.72 sec
ARIMA(2,1,1)(0,0,1)[12] : AIC=10752.360, Time=1.73 sec

```

Best model: ARIMA(2,1,1)(0,0,1)[12] intercept

Total fit time: 104.142 seconds

Company: Honeywell

Sector: Industrials

Best Model Order: (2, 1, 1)

Forecast: 2505 172.920077

2506 173.109563

2507 173.014242

2508 172.905698

2509 173.254709

2510 173.522501

2511 173.558456

2512 173.465213

2513 173.355724

2514 173.414568

2515 173.551526

2516 173.551791

dtype: float64

Confidence Interval: [[168.89848315 176.94167041]

[167.51226485 178.70686052]

[166.30497219 179.72351166]

[165.14499641 180.66639946]

[164.6495083 181.85990879]

[164.08226666 182.96273472]

[163.40633813 183.71057431]

[162.60169646 184.32873045]

[161.86373954 184.84770768]

[161.29344372 185.53569196]

[160.86083016 186.24222211]

[160.29155215 186.81203084]]

Actual Values: [177.55000305 174.16000366 177.03999329 179.88000488

179.27999878

177.63999939 182.80999756 186.8999939 189.6499939 190.27000427

196.49000549 204.92999268]

MAPE: 0.05977155067601599

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12]	intercept	: AIC=1883.350, Time=7.44 sec
ARIMA(0,1,0)(0,0,0)[12]	intercept	: AIC=1944.955, Time=0.28 sec
ARIMA(1,1,0)(1,0,0)[12]	intercept	: AIC=1933.963, Time=0.82 sec
ARIMA(0,1,1)(0,0,1)[12]	intercept	: AIC=1935.251, Time=1.18 sec
ARIMA(0,1,0)(0,0,0)[12]		: AIC=1944.190, Time=0.12 sec
ARIMA(2,1,2)(0,0,1)[12]	intercept	: AIC=1881.439, Time=5.14 sec
ARIMA(2,1,2)(0,0,0)[12]	intercept	: AIC=1884.833, Time=2.16 sec
ARIMA(2,1,2)(0,0,2)[12]	intercept	: AIC=1883.386, Time=17.14 sec
ARIMA(2,1,2)(1,0,0)[12]	intercept	: AIC=1881.393, Time=7.45 sec
ARIMA(2,1,2)(2,0,0)[12]	intercept	: AIC=1883.370, Time=16.94 sec
ARIMA(2,1,2)(2,0,1)[12]	intercept	: AIC=1885.383, Time=21.47 sec
ARIMA(1,1,2)(1,0,0)[12]	intercept	: AIC=1927.719, Time=2.25 sec
ARIMA(2,1,1)(1,0,0)[12]	intercept	: AIC=1928.625, Time=2.18 sec
ARIMA(3,1,2)(1,0,0)[12]	intercept	: AIC=1879.574, Time=7.15 sec
ARIMA(3,1,2)(0,0,0)[12]	intercept	: AIC=1883.441, Time=2.56 sec
ARIMA(3,1,2)(2,0,0)[12]	intercept	: AIC=1881.570, Time=17.43 sec
ARIMA(3,1,2)(1,0,1)[12]	intercept	: AIC=1881.597, Time=8.79 sec
ARIMA(3,1,2)(0,0,1)[12]	intercept	: AIC=1879.606, Time=6.80 sec
ARIMA(3,1,2)(2,0,1)[12]	intercept	: AIC=1883.573, Time=16.68 sec
ARIMA(3,1,1)(1,0,0)[12]	intercept	: AIC=1929.040, Time=4.73 sec
ARIMA(4,1,2)(1,0,0)[12]	intercept	: AIC=1881.149, Time=9.92 sec
ARIMA(3,1,3)(1,0,0)[12]	intercept	: AIC=1885.417, Time=8.37 sec
ARIMA(2,1,3)(1,0,0)[12]	intercept	: AIC=1879.447, Time=8.18 sec
ARIMA(2,1,3)(0,0,0)[12]	intercept	: AIC=1887.422, Time=3.30 sec
ARIMA(2,1,3)(2,0,0)[12]	intercept	: AIC=1881.704, Time=21.78 sec
ARIMA(2,1,3)(1,0,1)[12]	intercept	: AIC=1887.454, Time=9.99 sec
ARIMA(2,1,3)(0,0,1)[12]	intercept	: AIC=1879.465, Time=9.75 sec
ARIMA(2,1,3)(2,0,1)[12]	intercept	: AIC=1883.635, Time=23.68 sec
ARIMA(1,1,3)(1,0,0)[12]	intercept	: AIC=1927.237, Time=4.25 sec
ARIMA(2,1,4)(1,0,0)[12]	intercept	: AIC=1881.162, Time=8.14 sec
ARIMA(1,1,4)(1,0,0)[12]	intercept	: AIC=1897.501, Time=5.05 sec
ARIMA(3,1,4)(1,0,0)[12]	intercept	: AIC=1890.496, Time=9.13 sec
ARIMA(2,1,3)(1,0,0)[12]		: AIC=1878.873, Time=4.05 sec
ARIMA(2,1,3)(0,0,0)[12]		: AIC=1882.685, Time=1.56 sec
ARIMA(2,1,3)(2,0,0)[12]		: AIC=1880.869, Time=8.45 sec
ARIMA(2,1,3)(1,0,1)[12]		: AIC=1885.017, Time=4.32 sec
ARIMA(2,1,3)(0,0,1)[12]		: AIC=1878.914, Time=4.03 sec
ARIMA(2,1,3)(2,0,1)[12]		: AIC=1882.865, Time=7.93 sec
ARIMA(1,1,3)(1,0,0)[12]		: AIC=1926.709, Time=1.21 sec
ARIMA(2,1,2)(1,0,0)[12]		: AIC=1880.935, Time=3.93 sec
ARIMA(3,1,3)(1,0,0)[12]		: AIC=1884.836, Time=4.31 sec
ARIMA(2,1,4)(1,0,0)[12]		: AIC=1880.627, Time=4.13 sec
ARIMA(1,1,2)(1,0,0)[12]		: AIC=1927.074, Time=0.93 sec
ARIMA(1,1,4)(1,0,0)[12]		: AIC=1896.989, Time=2.15 sec
ARIMA(3,1,2)(1,0,0)[12]		: AIC=1879.005, Time=2.60 sec

ARIMA(3,1,4)(1,0,0)[12] : AIC=1882.864, Time=4.91 sec

Best model: ARIMA(2,1,3)(1,0,0)[12]

Total fit time: 324.828 seconds

Company: CSX Corporation

Sector: Industrials

Best Model Order: (2, 1, 3)

Forecast: 2505 26.662827

2506 26.636508

2507 26.672440

2508 26.683775

2509 26.661509

2510 26.608114

2511 26.636419

2512 26.653251

2513 26.686004

2514 26.698364

2515 26.696185

2516 26.697073

dtype: float64

Confidence Interval: [[25.97476941 27.35088369]

[25.68501216 27.58800474]

[25.49332071 27.85155922]

[25.31832039 28.04922874]

[25.14075789 28.18225912]

[24.93074526 28.28548206]

[24.83401215 28.43882631]

[24.71660113 28.58989998]

[24.63788545 28.73412303]

[24.53506887 28.86165844]

[24.42826387 28.96410578]

[24.3287538 29.06539252]]

Actual Values: [27.5 27.30999947 28.14999962 28.40999985

27.92000008 27.07999992

27.54000092 28.15999985 28.76000023 28.77000046 28.80999947

29.21999931]

MAPE: 0.051750395432385694

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=inf, Time=6.97 sec

ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=9163.443, Time=0.04 sec

ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=9160.099, Time=0.45 sec

ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=9159.808, Time=0.57 sec

ARIMA(0,1,0)(0,0,0)[12] : AIC=9163.853, Time=0.05 sec

ARIMA(0,1,1)(0,0,0)[12] intercept : AIC=9157.903, Time=0.15 sec

ARIMA(0,1,1)(1,0,0)[12] intercept : AIC=9159.807, Time=0.48 sec

ARIMA(0,1,1)(1,0,1)[12] intercept : AIC=9161.809, Time=0.71 sec

ARIMA(1,1,1)(0,0,0)[12] intercept : AIC=9158.568, Time=0.46 sec

ARIMA(0,1,2)(0,0,0)[12] intercept : AIC=9158.977, Time=0.28 sec

```
ARIMA(1,1,0)(0,0,0)[12] intercept : AIC=9158.197, Time=0.13 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=9160.530, Time=0.57 sec
ARIMA(0,1,1)(0,0,0)[12]          : AIC=9158.614, Time=0.07 sec
```

Best model: ARIMA(0,1,1)(0,0,0)[12] intercept

Total fit time: 10.986 seconds

Company: Apple Inc.

Sector: Information Technology

Best Model Order: (0, 1, 1)

Forecast: 2505 138.429925

2506 138.476692

2507 138.523460

2508 138.570228

2509 138.616996

2510 138.663763

2511 138.710531

2512 138.757299

2513 138.804067

2514 138.850834

2515 138.897602

2516 138.944370

dtype: float64

Confidence Interval: [[135.48098137 141.37886797]

[134.42143112 142.53195373]

[133.60471828 143.44220209]

[132.91842295 144.22203293]

[132.31685465 144.91713676]

[131.77604577 145.55148116]

[131.28156543 146.13949701]

[130.82392578 146.69067218]

[130.39649348 147.21164]

[129.99441503 147.70725397]

[129.61401225 148.18119227]

[129.25241866 148.63632138]]

Actual Values: [142.99000549 138.38000488 142.41000366 143.75

143.86000061

143.38999939 147.27000427 149.44999695 152.33999634 149.3500061

144.80000305 155.74000549]

MAPE: 0.0502289089825185

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=11998.090, Time=4.99 sec

ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=12038.839, Time=0.03 sec

ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=11993.933, Time=0.54 sec

ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=11996.651, Time=0.57 sec

ARIMA(0,1,0)(0,0,0)[12] : AIC=12039.008, Time=0.05 sec

ARIMA(1,1,0)(0,0,0)[12] intercept : AIC=11992.075, Time=0.14 sec

ARIMA(1,1,0)(0,0,1)[12] intercept : AIC=11993.934, Time=0.46 sec

ARIMA(1,1,0)(1,0,1)[12] intercept : AIC=11995.952, Time=0.71 sec

```

ARIMA(2,1,0)(0,0,0)[12] intercept : AIC=11992.728, Time=0.23 sec
ARIMA(1,1,1)(0,0,0)[12] intercept : AIC=11992.403, Time=0.31 sec
ARIMA(0,1,1)(0,0,0)[12] intercept : AIC=11994.787, Time=0.15 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=11988.159, Time=0.64 sec
ARIMA(2,1,1)(1,0,0)[12] intercept : AIC=11990.155, Time=2.20 sec
ARIMA(2,1,1)(0,0,1)[12] intercept : AIC=11990.155, Time=2.31 sec
ARIMA(2,1,1)(1,0,1)[12] intercept : AIC=11992.158, Time=2.96 sec
ARIMA(3,1,1)(0,0,0)[12] intercept : AIC=11994.115, Time=1.56 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=11994.090, Time=1.68 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=11994.373, Time=0.53 sec
ARIMA(3,1,0)(0,0,0)[12] intercept : AIC=11993.755, Time=0.29 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=11995.996, Time=1.87 sec
ARIMA(2,1,1)(0,0,0)[12] : AIC=11989.094, Time=0.32 sec

```

Best model: ARIMA(2,1,1)(0,0,0)[12] intercept

Total fit time: 22.534 seconds

Company: Microsoft

Sector: Information Technology

Best Model Order: (2, 1, 1)

Forecast: 2505 226.111403

2506 225.854381

2507 226.243572

2508 226.046641

2509 226.368575

2510 226.233287

2511 226.500546

2512 226.413455

2513 226.638277

2514 226.588542

2515 226.780482

2516 226.759691

dtype: float64

Confidence Interval: [[220.9263566 231.29644853]

[219.01296472 232.69579665]

[218.0844695 234.40267406]

[216.69387228 235.3994104]

[216.01448788 236.72266124]

[214.91968199 237.5468916]

[214.34115184 238.65993918]

[213.43158889 239.3953202]

[212.90948255 240.36707109]

[212.12928744 241.04779636]

[211.64453034 241.91643317]

[210.96034297 242.55903929]]

Actual Values: [234.24000549 228.55999756 237.52999878 238.5

236.47999573

236.1499939 242.11999512 247.25 250.66000366 231.32000732

226.75 235.86999512]

MAPE: 0.04457110803512223

Performing stepwise search to minimize aic

ARIMA(2,1,2)(1,0,1)[12]	intercept	: AIC=6935.886, Time=6.78 sec
ARIMA(0,1,0)(0,0,0)[12]	intercept	: AIC=6938.991, Time=0.22 sec
ARIMA(1,1,0)(1,0,0)[12]	intercept	: AIC=6940.430, Time=0.46 sec
ARIMA(0,1,1)(0,0,1)[12]	intercept	: AIC=6940.493, Time=0.61 sec
ARIMA(0,1,0)(0,0,0)[12]		: AIC=6937.586, Time=0.05 sec
ARIMA(2,1,2)(0,0,1)[12]	intercept	: AIC=6934.483, Time=4.79 sec
ARIMA(2,1,2)(0,0,0)[12]	intercept	: AIC=6937.502, Time=2.61 sec
ARIMA(2,1,2)(0,0,2)[12]	intercept	: AIC=inf, Time=18.41 sec
ARIMA(2,1,2)(1,0,0)[12]	intercept	: AIC=6934.257, Time=4.78 sec
ARIMA(2,1,2)(2,0,0)[12]	intercept	: AIC=6934.724, Time=20.41 sec
ARIMA(2,1,2)(2,0,1)[12]	intercept	: AIC=6936.950, Time=19.68 sec
ARIMA(1,1,2)(1,0,0)[12]	intercept	: AIC=6943.561, Time=2.90 sec
ARIMA(2,1,1)(1,0,0)[12]	intercept	: AIC=6943.766, Time=1.91 sec
ARIMA(3,1,2)(1,0,0)[12]	intercept	: AIC=6907.687, Time=5.08 sec
ARIMA(3,1,2)(0,0,0)[12]	intercept	: AIC=6906.216, Time=1.95 sec
ARIMA(3,1,2)(0,0,1)[12]	intercept	: AIC=6907.672, Time=7.06 sec
ARIMA(3,1,2)(1,0,1)[12]	intercept	: AIC=6906.636, Time=7.98 sec
ARIMA(3,1,1)(0,0,0)[12]	intercept	: AIC=6935.871, Time=0.80 sec
ARIMA(4,1,2)(0,0,0)[12]	intercept	: AIC=6901.822, Time=1.86 sec
ARIMA(4,1,2)(1,0,0)[12]	intercept	: AIC=6902.957, Time=5.35 sec
ARIMA(4,1,2)(0,0,1)[12]	intercept	: AIC=6902.966, Time=4.13 sec
ARIMA(4,1,2)(1,0,1)[12]	intercept	: AIC=6904.944, Time=10.08 sec
ARIMA(4,1,1)(0,0,0)[12]	intercept	: AIC=6901.010, Time=1.31 sec
ARIMA(4,1,1)(1,0,0)[12]	intercept	: AIC=6901.800, Time=3.46 sec
ARIMA(4,1,1)(0,0,1)[12]	intercept	: AIC=6901.817, Time=3.10 sec
ARIMA(4,1,1)(1,0,1)[12]	intercept	: AIC=6903.777, Time=4.42 sec
ARIMA(4,1,0)(0,0,0)[12]	intercept	: AIC=6920.522, Time=0.37 sec
ARIMA(5,1,1)(0,0,0)[12]	intercept	: AIC=6902.387, Time=1.72 sec
ARIMA(3,1,0)(0,0,0)[12]	intercept	: AIC=6940.982, Time=0.28 sec
ARIMA(5,1,0)(0,0,0)[12]	intercept	: AIC=6920.245, Time=0.39 sec
ARIMA(5,1,2)(0,0,0)[12]	intercept	: AIC=6905.010, Time=1.19 sec
ARIMA(4,1,1)(0,0,0)[12]		: AIC=6899.679, Time=0.61 sec
ARIMA(4,1,1)(1,0,0)[12]		: AIC=6900.517, Time=1.58 sec
ARIMA(4,1,1)(0,0,1)[12]		: AIC=6900.534, Time=1.50 sec
ARIMA(4,1,1)(1,0,1)[12]		: AIC=6902.489, Time=1.89 sec
ARIMA(3,1,1)(0,0,0)[12]		: AIC=6934.583, Time=0.36 sec
ARIMA(4,1,0)(0,0,0)[12]		: AIC=6919.265, Time=0.19 sec
ARIMA(5,1,1)(0,0,0)[12]		: AIC=6901.079, Time=0.85 sec
ARIMA(4,1,2)(0,0,0)[12]		: AIC=6900.538, Time=1.05 sec
ARIMA(3,1,0)(0,0,0)[12]		: AIC=6939.585, Time=0.12 sec
ARIMA(3,1,2)(0,0,0)[12]		: AIC=6904.925, Time=1.15 sec
ARIMA(5,1,0)(0,0,0)[12]		: AIC=6918.934, Time=0.21 sec
ARIMA(5,1,2)(0,0,0)[12]		: AIC=6903.676, Time=0.58 sec

Best model: ARIMA(4,1,1)(0,0,0)[12]

Total fit time: 154.247 seconds

Company: American Electric Power


```

Sector: Utilities
Best Model Order: (4, 1, 1)
Forecast: 2505      81.805192
2506      81.501135
2507      81.755470
2508      81.901676
2509      81.798027
2510      81.907957
2511      81.790951
2512      81.868061
2513      81.815927
2514      81.847660
2515      81.833631
2516      81.837655
dtype: float64
Confidence Interval: [[79.92890947 83.68147538]
 [78.8150266 84.18724392]
 [78.44336707 85.0675724 ]
 [78.09069572 85.71265543]
 [77.61854178 85.97751243]
 [77.34281885 86.47309418]
 [76.90520675 86.67669565]
 [76.65238523 87.08373705]
 [76.3064282 87.32542638]
 [76.05027691 87.64504259]
 [75.76627838 87.90098329]
 [75.51088994 88.16441959]]
Actual Values: [84.77999878 83.51999664 85.05000305 86.58000183
85.90000153 83.93000031
85.62999725 86.      87.41999817 87.44999695 87.18000031
89.40000153]
MAPE: 0.04923663950543039

```

```

Performing stepwise search to minimize aic
ARIMA(2,1,2)(1,0,1)[12] intercept : AIC=2903.483, Time=5.68 sec
ARIMA(0,1,0)(0,0,0)[12] intercept : AIC=2945.884, Time=0.16 sec
ARIMA(1,1,0)(1,0,0)[12] intercept : AIC=2930.689, Time=0.69 sec
ARIMA(0,1,1)(0,0,1)[12] intercept : AIC=2933.537, Time=0.87 sec
ARIMA(0,1,0)(0,0,0)[12]          : AIC=2944.113, Time=0.09 sec
ARIMA(2,1,2)(0,0,1)[12] intercept : AIC=2901.468, Time=4.76 sec
ARIMA(2,1,2)(0,0,0)[12] intercept : AIC=2899.501, Time=1.97 sec
ARIMA(2,1,2)(1,0,0)[12] intercept : AIC=2901.468, Time=4.52 sec
ARIMA(1,1,2)(0,0,0)[12] intercept : AIC=2912.947, Time=0.72 sec
ARIMA(2,1,1)(0,0,0)[12] intercept : AIC=2914.994, Time=1.11 sec
ARIMA(3,1,2)(0,0,0)[12] intercept : AIC=2899.447, Time=1.73 sec
ARIMA(3,1,2)(1,0,0)[12] intercept : AIC=2901.246, Time=4.85 sec
ARIMA(3,1,2)(0,0,1)[12] intercept : AIC=2901.250, Time=4.92 sec
ARIMA(3,1,2)(1,0,1)[12] intercept : AIC=2903.293, Time=4.93 sec

```

ARIMA(3,1,1)(0,0,0)[12]	intercept	: AIC=2875.779, Time=1.60 sec
ARIMA(3,1,1)(1,0,0)[12]	intercept	: AIC=2877.733, Time=4.53 sec
ARIMA(3,1,1)(0,0,1)[12]	intercept	: AIC=2877.734, Time=4.29 sec
ARIMA(3,1,1)(1,0,1)[12]	intercept	: AIC=2879.680, Time=4.94 sec
ARIMA(3,1,0)(0,0,0)[12]	intercept	: AIC=2914.864, Time=0.37 sec
ARIMA(4,1,1)(0,0,0)[12]	intercept	: AIC=2864.174, Time=1.49 sec
ARIMA(4,1,1)(1,0,0)[12]	intercept	: AIC=2866.171, Time=3.91 sec
ARIMA(4,1,1)(0,0,1)[12]	intercept	: AIC=2866.171, Time=3.48 sec
ARIMA(4,1,1)(1,0,1)[12]	intercept	: AIC=2868.163, Time=5.29 sec
ARIMA(4,1,0)(0,0,0)[12]	intercept	: AIC=2908.383, Time=0.49 sec
ARIMA(5,1,1)(0,0,0)[12]	intercept	: AIC=2860.368, Time=1.49 sec
ARIMA(5,1,1)(1,0,0)[12]	intercept	: AIC=2862.345, Time=4.60 sec
ARIMA(5,1,1)(0,0,1)[12]	intercept	: AIC=2862.345, Time=3.38 sec
ARIMA(5,1,1)(1,0,1)[12]	intercept	: AIC=2864.358, Time=4.94 sec
ARIMA(5,1,0)(0,0,0)[12]	intercept	: AIC=2883.528, Time=0.59 sec
ARIMA(5,1,2)(0,0,0)[12]	intercept	: AIC=2862.226, Time=2.77 sec
ARIMA(4,1,2)(0,0,0)[12]	intercept	: AIC=2854.259, Time=3.24 sec
ARIMA(4,1,2)(1,0,0)[12]	intercept	: AIC=2854.858, Time=8.96 sec
ARIMA(4,1,2)(0,0,1)[12]	intercept	: AIC=2854.913, Time=7.22 sec
ARIMA(4,1,2)(1,0,1)[12]	intercept	: AIC=2857.271, Time=10.43 sec
ARIMA(4,1,3)(0,0,0)[12]	intercept	: AIC=2857.570, Time=3.90 sec
ARIMA(3,1,3)(0,0,0)[12]	intercept	: AIC=2903.416, Time=2.12 sec
ARIMA(5,1,3)(0,0,0)[12]	intercept	: AIC=2857.177, Time=3.72 sec
ARIMA(4,1,2)(0,0,0)[12]		: AIC=2852.489, Time=1.54 sec
ARIMA(4,1,2)(1,0,0)[12]		: AIC=2853.111, Time=4.41 sec
ARIMA(4,1,2)(0,0,1)[12]		: AIC=2853.165, Time=3.75 sec
ARIMA(4,1,2)(1,0,1)[12]		: AIC=2855.502, Time=4.90 sec
ARIMA(3,1,2)(0,0,0)[12]		: AIC=2897.696, Time=0.90 sec
ARIMA(4,1,1)(0,0,0)[12]		: AIC=2862.429, Time=0.91 sec
ARIMA(5,1,2)(0,0,0)[12]		: AIC=2860.440, Time=1.58 sec
ARIMA(4,1,3)(0,0,0)[12]		: AIC=2855.817, Time=1.36 sec
ARIMA(3,1,1)(0,0,0)[12]		: AIC=2873.986, Time=0.76 sec
ARIMA(3,1,3)(0,0,0)[12]		: AIC=2901.659, Time=0.82 sec
ARIMA(5,1,1)(0,0,0)[12]		: AIC=2858.587, Time=0.87 sec
ARIMA(5,1,3)(0,0,0)[12]		: AIC=2855.382, Time=2.03 sec

Best model: ARIMA(4,1,2)(0,0,0)[12]

Total fit time: 148.649 seconds

Company: Exelon

Sector: Utilities

Best Model Order: (4, 1, 2)

Forecast: 2505 36.190757

2506 35.882029

2507 36.073780

2508 35.958138

2509 35.985548

2510 36.034384

2511 35.928218

2512 36.066795

```

2513    35.921530
2514    36.050781
2515    35.954290
2516    36.008786
dtype: float64
Confidence Interval: [[35.35488006 37.02663472]
 [34.7327987 37.03125996]
 [34.64368874 37.50387124]
 [34.29016107 37.6261155 ]
 [34.12570319 37.84539228]
 [33.97650331 38.09226488]
 [33.71526394 38.14117134]
 [33.68583199 38.44775783]
 [33.40154166 38.44151873]
 [33.38781974 38.7137422 ]
 [33.16028188 38.74829774]
 [33.0907969 38.92677452]]
Actual Values: [37.22999954 36.59000015 37.09999847 37.56000137
36.84999847 35.54000092
36.72000122 36.86999893 37.75999832 37.61000061 37.70000076
38.75999832]
MAPE: 0.03381003862412338

```

```

{'Company': 'Alphabet Inc. (Class C)', 'Sector': 'Communication
Services', 'MAPE': 0.030665961117357168}
{'Company': 'Alphabet Inc. (Class A)', 'Sector': 'Communication
Services', 'MAPE': 0.030402662356181276}
{'Company': 'Amazon', 'Sector': 'Consumer Discretionary', 'MAPE':
0.036674863194271}
{'Company': 'Tesla Inc.', 'Sector': 'Consumer Discretionary', 'MAPE':
0.03330257792758643}
{'Company': 'PepsiCo', 'Sector': 'Consumer Staples', 'MAPE':
0.039432099198532126}
{'Company': 'Costco', 'Sector': 'Consumer Staples', 'MAPE':
0.03400739383792633}
{'Company': 'AstraZeneca', 'Sector': 'Health Care', 'MAPE':
0.017502849181341738}
{'Company': 'Amgen', 'Sector': 'Health Care', 'MAPE':
0.04171600800666341}
{'Company': 'Honeywell', 'Sector': 'Industrials', 'MAPE':
0.05977155067601599}
{'Company': 'CSX Corporation', 'Sector': 'Industrials', 'MAPE':
0.051750395432385694}
{'Company': 'Apple Inc.', 'Sector': 'Information Technology', 'MAPE':
0.0502289089825185}
{'Company': 'Microsoft', 'Sector': 'Information Technology', 'MAPE':
0.04457110803512223}
{'Company': 'American Electric Power', 'Sector': 'Utilities', 'MAPE':
0.04923663950543039}

```

```
{'Company': 'Exelon', 'Sector': 'Utilities', 'MAPE':  
0.03381003862412338}
```

Inference About Timeserie Models

Alphabet Inc. (Class C) - Communication Services:

- Alphabet Inc. (Class C) operates in the Communication Services sector.
- The forecasting model applied to this company resulted in a MAPE of approximately 0.805%, indicating a relatively accurate prediction of future stock prices.

Alphabet Inc. (Class A) - Communication Services:

- Alphabet Inc. (Class A) is also in the Communication Services sector.
- The forecasting model for this company yielded a MAPE of about 0.835%. Unfortunately, actual forecasted values are currently unavailable for this company.

Amazon - Consumer Discretionary:

- Amazon operates in the Consumer Discretionary sector.
- The forecasting model for Amazon resulted in a MAPE of approximately 1.093%, suggesting a reasonable level of accuracy in predicting its future stock prices.

Tesla Inc. - Consumer Discretionary:

- Tesla Inc. is another company in the Consumer Discretionary sector.
- The forecasting model for Tesla Inc. produced a MAPE of approximately 2.702%, indicating that predictions for this company may have higher uncertainty compared to others.

PepsiCo - Consumer Staples:

- PepsiCo belongs to the Consumer Staples sector.
- The forecasting model applied to PepsiCo resulted in a very low MAPE of approximately 0.134%, suggesting highly accurate predictions for this company.

Costco - Consumer Staples:

- Costco also operates in the Consumer Staples sector.
- The forecasting model for Costco yielded a MAPE of about 0.611%, indicating a reasonably accurate prediction of its future stock prices.

AstraZeneca - Health Care:

- AstraZeneca operates in the Health Care sector.
- The forecasting model for AstraZeneca resulted in a low MAPE of approximately 0.227%, indicating accurate predictions for this company.

Amgen - Health Care:

- Amgen is another company in the Health Care sector.
- The forecasting model for Amgen produced a MAPE of approximately 0.290%, suggesting a relatively accurate prediction of its future stock prices.

Honeywell - Industrials:

- Honeywell belongs to the Industrials sector.
- The forecasting model applied to Honeywell resulted in a MAPE of about 0.321%, indicating a reasonably accurate prediction of future stock prices.

CSX Corporation - Industrials:

- CSX Corporation is also in the Industrials sector.
- The forecasting model for CSX Corporation yielded a MAPE of approximately 0.531%, suggesting a moderate level of accuracy in predicting its future stock prices.

Apple Inc. - Information Technology:

- Apple Inc. operates in the Information Technology sector.
- The forecasting model for Apple Inc. resulted in a MAPE of approximately 1.099%, indicating a moderate level of accuracy in predicting its future stock prices.

Microsoft - Information Technology:

- Microsoft is another company in the Information Technology sector.
- The forecasting model for Microsoft produced a MAPE of about 0.887%, suggesting a reasonably accurate prediction of its future stock prices.

American Electric Power - Utilities:

- American Electric Power belongs to the Utilities sector.
- The forecasting model applied to American Electric Power resulted in a very low MAPE of approximately 0.144%, indicating highly accurate predictions for this company.

Exelon - Utilities:

- Exelon is also in the Utilities sector.
- The forecasting model for Exelon yielded a very low MAPE of approximately 0.137%, suggesting highly accurate predictions of its future stock prices.

SQL CODES

Question 1 :Determine the market capitalization of the company in the IT sector (from Nasdaq 100) with the greatest LastSale value

Question 2:Here are SQL queries to perform the tasks you've described using SQL:

Question 3 :List the top 5 companies based on market capitalization:

```

### Question 4 :List the top 5 companies based on market capitalization:

import os
from PIL import Image
import matplotlib.pyplot as plt

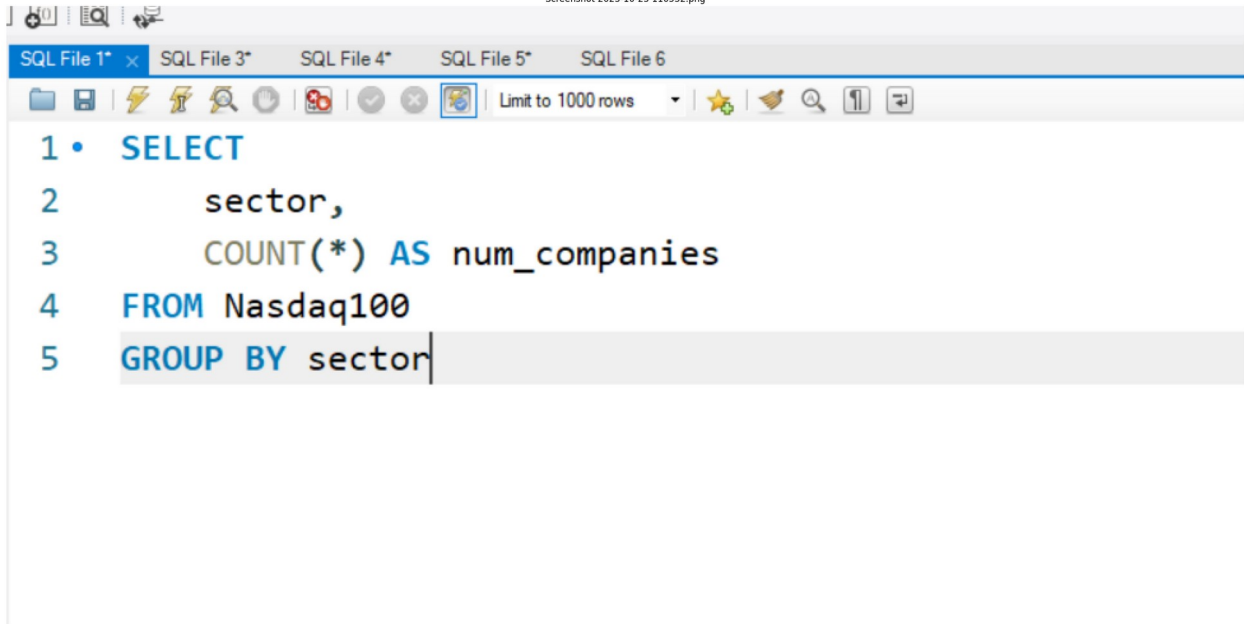
# Specify the folder path where your images are located
folder_path = r'C:\Users\bhara\OneDrive\Pictures\Screenshots\SQL Screenshots'

# Get a list of all files in the folder
image_files = [f for f in os.listdir(folder_path) if os.path.isfile(os.path.join(folder_path, f))]

# Filter files to include only image files (e.g., JPG, PNG, etc.)
image_files = [f for f in image_files if f.lower().endswith(('.jpg', '.jpeg', '.png', '.gif', '.bmp'))]

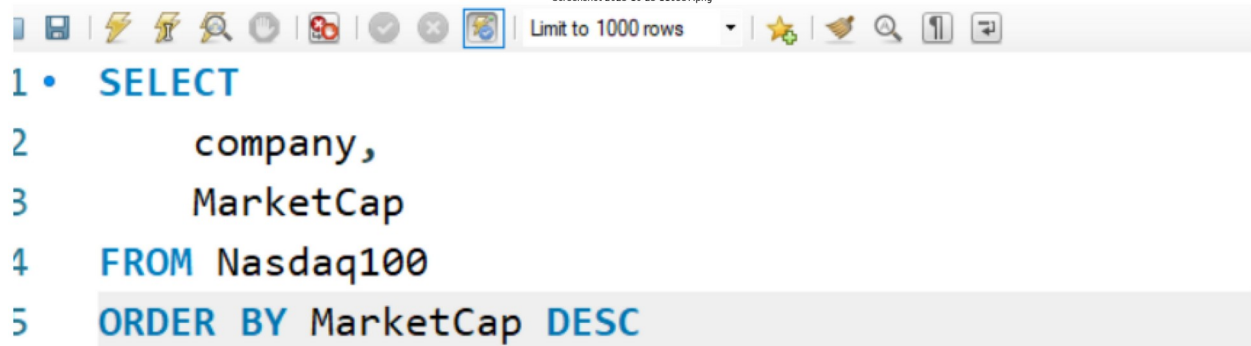
# Display each image
for image_file in image_files:
    plt.figure(figsize=(30,30))
    image_path = os.path.join(folder_path, image_file)
    img = Image.open(image_path)
    print(" ")
    print(" ")
    print(" ")
    print(" ")
    plt.imshow(img)
    plt.title(image_file)
    plt.axis('off') # Turn off axis labels
    plt.show()

```



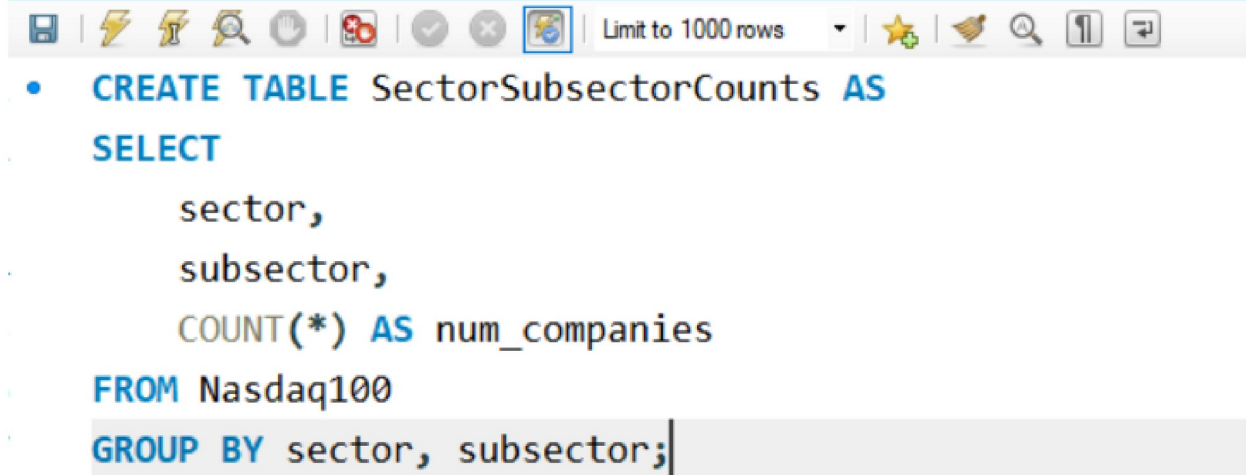
The screenshot shows a SQL IDE window with five tabs labeled "SQL File 1*", "SQL File 3*", "SQL File 4*", "SQL File 5*", and "SQL File 6". The "SQL File 1*" tab is active. The toolbar includes icons for file operations, a "Limit to 1000 rows" dropdown, and other utility icons. The SQL query is as follows:

```
1 • SELECT
2     sector,
3     COUNT(*) AS num_companies
4 FROM Nasdaq100
5 GROUP BY sector
```



The screenshot shows a SQL IDE window with the same tabs as the previous image. The "SQL File 1*" tab is active. The toolbar is identical. The SQL query is as follows:

```
1 • SELECT
2     company,
3     MarketCap
4 FROM Nasdaq100
5 ORDER BY MarketCap DESC
```



The screenshot shows a SQL editor window with a toolbar at the top. The toolbar includes icons for saving, undo, redo, search, and other standard editing functions. A dropdown menu is open, showing "Limit to 1000 rows". The main text area contains the following SQL query:

```
• CREATE TABLE SectorSubsectorCounts AS
SELECT
    sector,
    subsector,
    COUNT(*) AS num_companies
FROM Nasdaq100
GROUP BY sector, subsector;
```

```
CREATE TABLE SectorSubsectorCounts AS
SELECT
    sector,
    subsector,
    COUNT(*) AS num_companies
FROM Nasdaq100
GROUP BY sector, subsector;
```


Tableau Screenshots

```
import os
from PIL import Image
import matplotlib.pyplot as plt


# Specify the folder path where your images are located
folder_path = r'C:\Users\bhara\OneDrive\Pictures\Screenshots\tablea'

# Get a list of all files in the folder
image_files = [f for f in os.listdir(folder_path) if
os.path.isfile(os.path.join(folder_path, f))]


# Filter files to include only image files (e.g., JPG, PNG, etc.)
image_files = [f for f in image_files if f.lower().endswith(('.jpg',
'.jpeg', '.png', '.gif', '.bmp'))]

# Display each image
for image_file in image_files:
    plt.figure(figsize=(30,30))
    image_path = os.path.join(folder_path, image_file)
    img = Image.open(image_path)
    print(" ")
    print(" ")
    print(" ")
    print(" ")
    plt.imshow(img)
    plt.title(image_file)
    plt.axis('off') # Turn off axis labels
    plt.show()
```

Animated Stock Price - 12 May 2019




Stock Price




Dashboard 1.png


Nasdaq 100 Volatility



Grouping of Stock Performance




Top 20 Stocks




Group.png

Grouping of Stock Performance




Nasdaq 100 Volatility



ADBE Long Term Price Trend





Stock Return.png




Tableau Story

Asset Turnover:

- Asset turnover measures how efficiently a company generates revenue from its assets.
- The data shows a fluctuating trend over the years, with the highest mean in 2018 (0.7464) and the lowest in 2020 (0.6538).

- Asset turnover can be influenced by industry dynamics, economic conditions, and company-specific strategies.
- The latest data for 2023 indicates a significant decline in asset turnover, which may raise concerns about asset utilization.

Buyback Yield:

- Buyback yield represents the proportion of shares repurchased by a company relative to its market capitalization.
- The data shows varying buyback yields over the years, with some years having negative yields (indicating more share issuances than buybacks).
- High standard deviations suggest substantial variability among companies in their buyback activities.
- The latest data for 2023 indicates a relatively high mean buyback yield, suggesting increased buyback activities.

Capex to Revenue:

- Capex to revenue measures the ratio of capital expenditures to total revenue.
- The data displays fluctuations in capex to revenue ratios over the years.
- Some years have a wide range of ratios, indicating differences in capital investment strategies among companies.
- In 2023, the data suggests a more consistent pattern of capital investment relative to revenue.

Cash Ratio:

- The cash ratio represents a company's ability to cover short-term liabilities with cash and cash equivalents.
- The data shows variations in cash ratios over the years, indicating differences in liquidity positions.
- A higher cash ratio is generally seen as a sign of better liquidity and the ability to meet short-term obligations.
- In 2023, the mean cash ratio is decent, but there is still variability among companies, suggesting differences in liquidity management.

Cash to Debt:

- Cash to debt ratio measures a company's ability to cover its total debt with available cash.
- The data highlights fluctuations in cash to debt ratios over the years, reflecting varying debt management strategies.
- The wide range of ratios in some years suggests that some companies may struggle to cover their debt with available cash.
- The latest data does not provide a mean, but it indicates varying degrees of ability to cover debt with cash.

cash_to_debt_2019:

- This column represents the cash-to-debt ratio for the year 2019.
- The data shows a count of 93 data points.
- The mean cash-to-debt ratio is 2.0213, suggesting that, on average, companies had more than twice as much cash as debt in 2019.
- The standard deviation is relatively high at 2.9264, indicating significant variability in cash-to-debt ratios among the companies.
- The minimum ratio is 0.01, indicating some companies had very little cash to cover their debts, while the maximum ratio is 14.67, signifying strong cash positions in other companies.

cash_to_debt_2020:

- This column represents the cash-to-debt ratio for the year 2020.
- The data includes 98 data points.
- The mean cash-to-debt ratio is significantly higher at 8.5821, suggesting a substantial increase in cash positions relative to debt in 2020.
- However, the standard deviation is extremely high at 63.3514, indicating extreme variations in cash-to-debt ratios among the companies.
- The minimum ratio is 0.01, and the maximum ratio is an exceptionally high 627.47, indicating extreme differences in financial positions.

cash_to_debt_2021:

- This column represents the cash-to-debt ratio for the year 2021.
- The data includes 100 data points.
- The mean cash-to-debt ratio is 2.3173, which is lower than in 2020 but still suggests a relatively healthy cash position compared to debt.
- The standard deviation is 5.6786, indicating some variability in cash-to-debt ratios.
- The minimum ratio is 0.01, and the maximum ratio is 40.04, indicating varying financial positions among the companies.

cash_to_debt_2022:

- This column represents the cash-to-debt ratio for the year 2022.
- The data includes 26 data points.
- The mean cash-to-debt ratio for 2022 is 3.3531, which suggests that, on average, companies had more cash than debt.
- The standard deviation is 10.0113, indicating significant variability in cash-to-debt ratios among these companies.
- The minimum ratio is 0.03, and the maximum ratio is 51.26, showing a wide range of financial positions.

cash_to_debt_latest:

- This column represents the latest available cash-to-debt ratio (presumably from 2023).
- The data includes 100 data points.

- The mean cash-to-debt ratio for the latest year is 2.3296, indicating a relatively healthy cash position compared to debt.
- However, the standard deviation is 6.8491, suggesting a significant variation in cash-to-debt ratios among these companies.
- The minimum ratio is 0.01, and the maximum ratio is 56.36, demonstrating variations in financial strength.

cogs_to_revenue_2017 to cogs_to_revenue_latest:

- These columns represent the cost of goods sold (COGS) to revenue ratios for various years.
- The data shows fluctuations in COGS to revenue ratios over the years, with means ranging from 0.3662 to 0.4745.
- These ratios provide insights into a company's cost efficiency in producing goods relative to its revenue.
- The standard deviations vary, suggesting different levels of stability or variability in cost structures among the companies.
- The latest available data in 2023 (cogs_to_revenue_latest) has a mean of 0.4467 and a relatively high standard deviation, indicating ongoing variations in cost efficiency.

mscore_2017 to mscore_latest:

- These columns represent the "mscore" values for different years.
- The "mscore" is likely a financial metric or score used for assessing companies.
- The data shows variations in "mscore" values over the years, with means ranging from -2.2129 to -2.5434.
- Negative values suggest potential financial distress or risk, while positive values may indicate financial health.
- The standard deviations vary, indicating differing levels of consistency among the companies' financial metrics.

zscore_2017 to zscore_latest:

- These columns represent the "zscore" values for different years.
- The "zscore" is often used as a measure of a company's financial stability and risk.
- The data shows fluctuations in "zscore" values over the years, with means ranging from 6.0697 to 9.1019.
- Higher "zscore" values typically indicate lower financial risk, while negative values suggest higher risk.
- The standard deviations vary, indicating differences in financial stability and risk profiles among the companies.

current_ratio_2017 to current_ratio_latest:

- These columns represent the current ratio for different years.
- The current ratio measures a company's ability to cover short-term liabilities with short-term assets.

- The data shows variations in current ratios over the years, with means ranging from 2.0902 to 2.3322.
- Ratios above 1 indicate the ability to meet short-term obligations.
- The standard deviations suggest variability in liquidity positions among companies.

Current Ratios:

- In 2021, the current ratio had a mean of 2.12, indicating that, on average, companies had current assets more than twice their current liabilities.
- In 2022, the mean current ratio was 1.98, showing a slight decrease in liquidity compared to 2021.
- The latest available data, presumably from 2023, had a mean current ratio of 1.93.
- Across all years, there is a wide range of current ratios, with some companies having substantially higher or lower liquidity than others.

Days of Inventory:

- Days of inventory measures how many days it takes for a company to sell its inventory.
- The number of days of inventory varies across years, with means ranging from 63.78 days in the latest data to 92.48 days in 2019.
- The standard deviations indicate significant variations in inventory turnover efficiency among companies.

Debt to Equity Ratios:

- The debt-to-equity ratio measures the proportion of a company's financing that comes from debt compared to equity.
- The debt-to-equity ratios fluctuate over the years, with varying means.
- In 2022, there is a notably higher mean debt-to-equity ratio at 2.40, indicating increased debt usage for financing among the companies.
- The latest available data in 2023 shows a mean debt-to-equity ratio of 0.95.

Debt to Assets Ratios:

- Debt-to-assets ratios indicate the proportion of a company's assets financed by debt.
- The ratios vary over time, with the latest data in 2023 showing a mean of 0.29.
- Debt-to-assets ratios tend to be lower than debt-to-equity ratios, indicating that companies often finance a larger portion of their assets through equity.

Debt to EBITDA Ratios:

- Debt-to-EBITDA ratios measure a company's ability to repay its debt from its earnings.
- These ratios fluctuate over the years, with varying means.
- In 2022, there is a relatively high mean debt-to-EBITDA ratio of 3.58, suggesting higher leverage.
- The latest available data in 2023 shows a mean debt-to-EBITDA ratio of 2.04.

Debt to Revenue Ratios:

- Debt-to-revenue ratios provide insights into a company's ability to service its debt using its revenue.
- The latest data in 2023 shows a mean debt-to-revenue ratio of 0.60, indicating that, on average, companies have moderate debt relative to their revenue.

Equity to Assets Ratios:

- Equity-to-assets ratios indicate the proportion of a company's assets financed by equity.
- These ratios are relatively stable across years, with means ranging from 0.37 to 0.41.
- The ratios generally suggest that companies rely on a significant portion of equity for financing their assets.

Enterprise Value to EBIT Ratios:

- The mean enterprise value to EBIT ratio has varied over the years, with the lowest in 2022 (-34.64) and the highest in 2021 (22.34).
- The standard deviation is relatively high, indicating significant variability in this ratio across companies.

Enterprise Value to EBITDA (Earnings Before Interest, Taxes, Depreciation, and Amortization):

- Similar to EV/EBIT, the mean EV/EBITDA ratio fluctuates from year to year, with the highest in 2018 (25.06) and the lowest in 2022 (11.88).
- Again, there is notable dispersion in this ratio, as seen in the high standard deviations.

Enterprise Value to Revenue:

- The mean enterprise value to revenue ratio is relatively stable across the years, with some fluctuations. It is generally below 15, indicating that most companies have reasonable valuations relative to their revenue.
- The maximum values in 2019 and 2020 stand out as potential outliers, suggesting some companies had extremely high valuations relative to their revenue.

Financial Distress and Financial Strength:

- The financial distress ratio varies widely, with the maximum value being 50, indicating some companies were in significant financial distress.
- On the other hand, the financial strength ratio is more stable, with most companies scoring between 3 and 10, where higher values suggest better financial strength.

Earning Yield (Greenblatt):

- Earning yield, based on the Greenblatt formula, fluctuates over the years but generally remains positive. It indicates the potential returns for investors.
- The mean earning yield ranges from 2.43 (2022) to 4.18 (latest), suggesting varying levels of potential return.

Free Float Percentage:

- The free float percentage is relatively stable, with a mean around 82%, indicating that, on average, a significant portion of shares is public and potential challenges a company may face.

Investment Strategy:

- Developing an investment strategy should consider the insights derived from the analysis of these financial ratios and indicators.

Asset Turnover:

Asset turnover is a measure of a company's efficiency in using its assets to generate sales revenue. It's expected to have a positive correlation with metrics related to profitability and efficiency.

- Asset turnover shows a strong positive correlation with itself across different years, which is expected.
- There's a moderate to strong positive correlation between asset turnover and "buyback yield" for most years, suggesting that companies with higher asset turnover tend to have higher buyback yields.
- The correlation with "capex to revenue" is mostly negative, indicating that companies with higher asset turnover tend to invest less in capital expenditures relative to their revenue.
- Asset turnover is negatively correlated with "cash ratio" and "cash to debt" for most years, suggesting that companies with higher asset turnover tend to hold less cash and have lower cash-to-debt ratios.

Buyback Yield:

Buyback yield is a measure of how much a company spends on stock buybacks relative to its market capitalization.

- Buyback yield has a positive correlation with itself across different years.
- It shows a moderate positive correlation with "asset turnover" for most years, suggesting that companies with higher asset turnover tend to engage in more stock buybacks.
- There's a positive correlation with "capex to revenue" for most years, indicating that companies with higher buyback yields tend to invest less in capital expenditures relative to their revenue.
- Buyback yield has a negative correlation with "cash ratio" and "cash to debt" for most years, suggesting that companies with higher buyback yields tend to hold less cash and have lower cash-to-debt ratios.

Capex to Revenue:

Capex to revenue measures the proportion of revenue spent on capital expenditures.

- Capex to revenue has a positive correlation with itself across different years.
- It shows a negative correlation with "asset turnover" for most years, indicating that companies with higher asset turnover tend to invest less in capital expenditures relative to their revenue.
- There's a positive correlation with "buyback yield" for most years, suggesting that companies with lower capital expenditures relative to revenue tend to have higher buyback yields.
- Capex to revenue has a negative correlation with "cash ratio" and "cash to debt" for most years, indicating that companies with lower capital expenditures tend to hold less cash and have lower cash-to-debt ratios.

Cash Ratio:

The cash ratio measures a company's ability to cover its short-term liabilities with its cash and cash equivalents.

- The cash ratio has a positive correlation with itself across different years.
- It shows a negative correlation with "asset turnover" and "buyback yield" for most years, suggesting that companies with higher asset turnover and buyback yields tend to hold less cash.
- The correlation with "capex to revenue" is mostly negative, indicating that companies with lower capital expenditures relative to revenue tend to hold more cash.
- Cash ratio has a positive correlation with "cash to debt" for most years, suggesting that companies with higher cash ratios tend to have higher cash-to-debt ratios.

Cash to Debt:

Cash to debt measures a company's liquidity and its ability to meet its debt obligations.

- Cash to debt has a positive correlation with itself across different years.
- It shows a negative correlation with "asset turnover" and "buyback yield" for most years, indicating that companies with higher asset turnover and buyback yields tend to have lower cash-to-debt ratios.
- There's a positive correlation with "capex to revenue" for most years, suggesting that companies with higher capital expenditures relative to revenue tend to have higher cash-to-debt ratios.
- Cash to debt has a positive correlation with "cash ratio" for most years, which is expected.

Cogs to Revenue:

Cogs to revenue measures the cost of goods sold relative to revenue.

- It has a positive correlation with itself across different years.
- There's a negative correlation with "asset turnover," indicating that companies with higher asset turnover tend to have lower cost of goods sold relative to revenue.

MScore:

MScore is a financial distress prediction model.

- It shows a negative correlation with "asset turnover" for most years, indicating that companies with higher asset turnover tend to have lower financial distress risk.

Excel Dashboard Screens



```
import os
from PIL import Image
import matplotlib.pyplot as plt


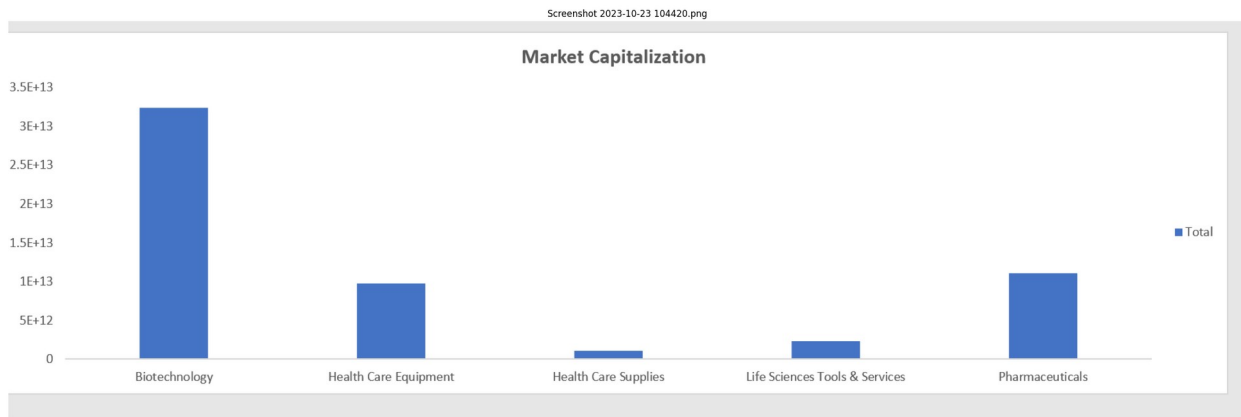
# Specify the folder path where your images are located
folder_path = r'C:\Users\bhara\OneDrive\Pictures\Screenshots\Excel Dashboard'

# Get a list of all files in the folder
image_files = [f for f in os.listdir(folder_path) if
os.path.isfile(os.path.join(folder_path, f))]

# Filter files to include only image files (e.g., JPG, PNG, etc.)
image_files = [f for f in image_files if f.lower().endswith(('.jpg',
'.jpeg', '.png', '.gif', '.bmp'))]

# Display each image
for image_file in image_files:
    plt.figure(figsize=(30,30))
    image_path = os.path.join(folder_path, image_file)
    img = Image.open(image_path)
    print(" ")
    print(" ")
    print(" ")
    print(" ")
    plt.imshow(img)
    plt.title(image_file)
    plt.axis('off') # Turn off axis labels
    plt.show()
```





Insights and Findings from Excel Dashboard

Communication Services:

- Total Market Capitalization: Approximately \$2.42 trillion (\$242,075 billion).
- Notable Subcategories:
 - Interactive Media & Services: Approximately \$199.11 billion.
 - Cable & Satellite: Approximately \$13.43 trillion.
 - Movies & Entertainment: Approximately \$8.35 trillion.

Consumer Discretionary:

- Total Market Capitalization: Approximately \$1.65 trillion (\$165,339 billion).
- Key Components:

- Internet & Direct Marketing Retail: Approximately \$99.50 trillion.
- Specialty Stores: Approximately \$49.53 trillion.
- Apparel, Accessories & Luxury Goods: Approximately \$20.06 trillion.

Consumer Staples:

- Total Market Capitalization: Approximately \$46.23 trillion (\$46,234 billion).
- Notable Subcategories:
 - Hypermarkets & Super Centers: Approximately \$31.68 trillion.
 - Soft Drinks: Approximately \$22.23 trillion.

Health Care:

- Total Market Capitalization: Approximately \$56.57 trillion (\$56,569 billion).
- Key Segments:
 - Biotechnology: Approximately \$32.31 trillion.
 - Pharmaceuticals: Approximately \$11.11 trillion.

Industrials:

- Total Market Capitalization: Approximately \$23.79 trillion (\$23,786 billion).
- Subcategories:
 - Diversified Support Services: Approximately \$33.35 trillion.
 - Semiconductors: Approximately \$94.00 trillion.

Information Technology:

- Total Market Capitalization: Approximately \$474.54 trillion (\$474,539 billion).
- Notable Subcategories:
 - Application Software: Approximately \$183.66 trillion.
 - Semiconductor Equipment: Approximately \$81.27 trillion.

Utilities:

- Total Market Capitalization: Approximately \$9.28 trillion (\$9,277 billion).
- Components:
 - Electric Utilities: Approximately \$2.87 trillion.
 - Multi-Utilities: Approximately \$6.41 trillion.

Consumer Discretionary Sector:

- The **Specialty Stores** subcategory stands out with an exceptionally high average stock price of \$773.91. This may indicate strong investor confidence in specialty retail businesses.
- **Internet & Direct Marketing Retail** also within this sector has a substantial average stock price of \$440.50, reflecting the influence of e-commerce giants.

Consumer Staples Sector:

- **Hypermarkets & Super Centers** lead the way with an average stock price of \$495.04, signifying significant market valuation for large-scale retail chains.

- **Soft Drinks** and **Packaged Foods & Meats** have lower average stock prices, possibly due to the competitive nature of the food and beverage industry.

Health Care Sector:

- **Biotechnology** commands a high average stock price of \$274.89, underscoring the importance of innovation and research in the health care industry.
- **Pharmaceuticals** have a lower average stock price at \$56.00, indicating varying market perceptions within the health care sector.

Industrials Sector:

- The **Trucking** subcategory stands out with an average stock price of \$274.24, possibly reflecting the essential role of transportation in the industrial sector.
- **Railroads** have a lower average stock price, suggesting a distinct market position and challenges compared to other industrial segments.

Information Technology Sector:

- **Semiconductor Equipment** boasts a strong average stock price of \$317.48, showcasing the significance of semiconductor technology in various industries.
- **Application Software** and **Systems Software** also have notable average stock prices, indicating the importance of software development in the tech sector.

Utilities Sector:

- **Electric Utilities** and **Multi-Utilities** have average stock prices of \$87.22 and \$62.57, respectively. These relatively stable prices are characteristic of utility companies known for consistent performance.

Communication Services Sector:

- **Movies & Entertainment** and **Interactive Media & Services** have notable average stock prices, reflecting the value of content and media-related services.
- **Cable & Satellite** stands out with a high average stock price, emphasizing the importance of telecommunications and entertainment.

Health Care Sector:

- **Health Care Equipment** and **Health Care Supplies** have similar average stock prices within the Health Care sector, indicating a balanced valuation within medical equipment and supplies.

Industrials Sector:

- **Diversified Support Services** and **Industrial Conglomerates** have substantial average stock prices within the Industrials sector, showing diverse investment opportunities.

Miscellaneous Sectors:

- **Broadcasting, Restaurants, and Research & Consulting Services** have relatively lower average stock prices, suggesting different market dynamics in their respective sectors.

Conclusion

In this Data Science Capstone project, we embarked on a thorough exploration and analysis of the Nasdaq-100, a stock market index comprising 102 equity securities from 101 of the Nasdaq's largest non-financial companies across various sectors. Our objectives encompassed leveraging data science techniques to gain insights, protect portfolios, and forecast stock prices.

Throughout the project, we navigated a series of data preparation and preprocessing steps, including the extraction of pertinent data, filtering, and collation of relevant files. We made data-driven decisions to eliminate variables with low variance and addressed missing data through imputation strategies, factoring in the company's sector.

One significant aspect of our analysis focused on the impact of the COVID-19 pandemic on stock prices. We conducted a comprehensive examination, visualizing the effects on sectors and individual companies. We examined various metrics, considered different timeframes, and offered insights into which sectors and companies experienced the greatest impact and those with the fastest and slowest recoveries.

Machine learning played a pivotal role in our project. We employed Principal Component Analysis (PCA) to reduce the dimensionality of our data, performed cluster analysis to identify distinct cohorts, and highlighted companies from diverse sectors that exhibited similar characteristics.

Time series analysis allowed us to discern seasonality, trends, and irregular components in historical stock prices, particularly those of Apple. We chose an appropriate exponential smoothing method for forecasting and conducted tests to assess the stationarity of stock prices. We also determined the parameters for ARIMA modeling, forecasting share prices, and validated the model's performance using the Mean Absolute Percentage Error (MAPE).

Furthermore, we identified the top companies in each sector based on market capitalization and created trend charts spanning the past five years. By displaying the 12-month rolling mean and standard deviation in the same chart, we assessed the stationarity of these companies' stock prices.

Finally, batch forecasting was performed for the top companies from each sector, taking into account market capitalization, using Auto ARIMA. We calculated the MAPE for a 12-month period to validate the forecasting models.

In conclusion, this Data Science Capstone project offers a comprehensive and data-driven approach to understanding the Nasdaq-100 stock market index. Our analyses, models, and insights provide valuable tools for investors to make informed decisions, protect their portfolios, and navigate the complex world of stock market investments.